


Chapter 4: Threads





Chapter 4: Threads

- ◆ Overview
- ◆ Multithreading Models
- ◆ Thread Libraries
- ◆ Examples(Java Threads, OSs Threads)
- ◆ Threading Issues





Objectives

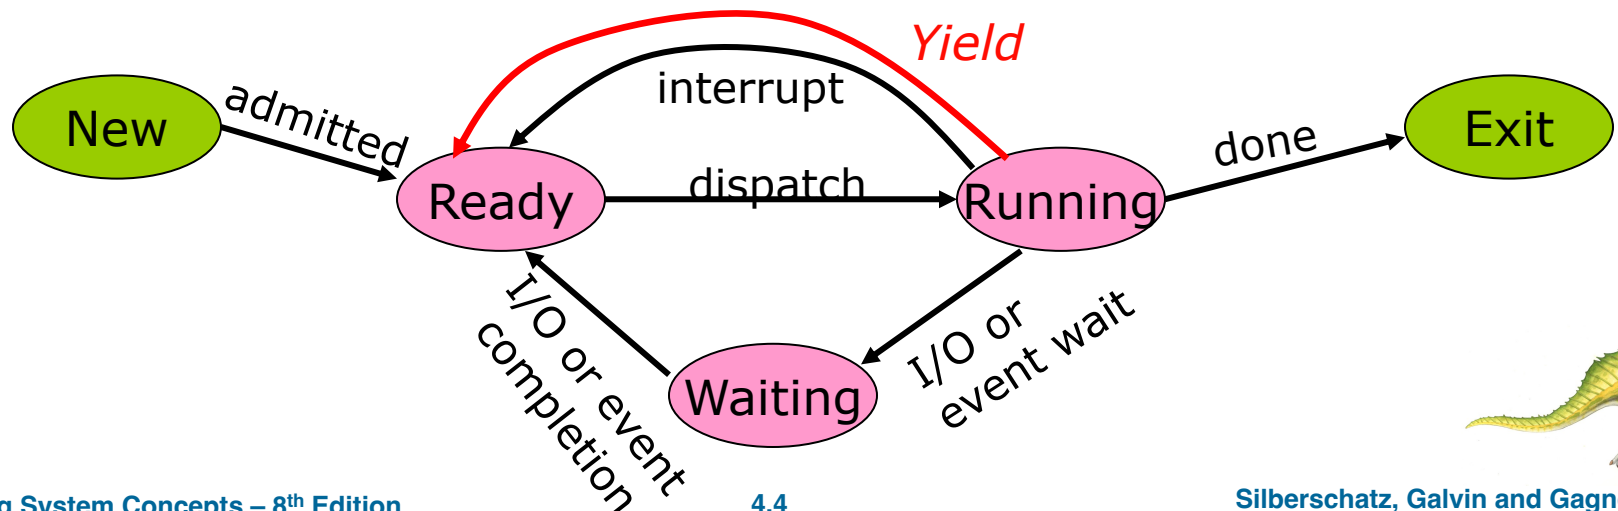
- ◆ To introduce the notion of a thread — a **fundamental unit of CPU utilization** that forms the basis of multithreaded computer systems
- ◆ To discuss the APIs for the Pthreads, Win32, and Java thread libraries
- ◆ To examine issues related to multithreaded programming





Review: Processes

- ◆ The unit of execution and scheduling
- ◆ A task created by the OS, running in a restricted virtual machine environment –a virtual CPU, virtual memory environment, interface to the OS via system calls
- ◆ Sequential, instruction-at-a-time execution of a program.
- ◆ Abstraction used for protection
 - Main Memory State (contents of Address Space)
- ◆ Multiprogramming: overlap IO and CPU
- ◆ Context Switches are expensive





Review: Processes

- ◆ How does parent know child process has finished if exec does not return?

```
main(int argc, char **argv)
{
    char *myName = argv[1];
    char *progName = argv[2];

    int cpid = fork();
    if (cpid == 0) {
        printf("The child of %s is %d\n", myName, getpid());
        execlp("/bin/ls", "ls", NULL);
        printf("OH NO. THEY LIED TO ME!!!\n");
    } else {
        printf("My child is %d\n", cpid);
        wait(cpid);
        exit(0);
    }
}
```





Review: Cooperating Processes

- ◆ Processes can be independent or work cooperatively
- ◆ Cooperating processes can be used:
 - to gain **speedup** by overlapping activities or working in parallel
 - to **better structure** an application as set of cooperating processes
 - to **share information** between jobs
- ◆ Sometimes processes are structured as a pipeline
 - each produces work for the next stage that consumes it

```
main()
  read_data()
  for(all data)
    compute();
    write_data();
  endfor
```

```
main()
  read_data()
  for(all data)
    compute();
    CreateProcess(write_data());
  endfor
```





Case for Parallelism

Consider the following code fragment on a dual core CPU

```
for(k = 0; k < n; k++)  
    a[k] = b[k] * c[k] + d[k] * e[k];
```

Instead:

```
CreateProcess(fn, 0, n/2);  
CreateProcess(fn, n/2, n);  
fn(l, m)  
    for(k = l; k < m; k++)  
        a[k] = b[k] * c[k] + d[k] * e[k];
```





Case for Parallelism

■ Consider a Web server:

create a number of process, and for each process do

- get network message from client
- get URL data from disk
- compose response
- send response

◆ Server connections are fast, but client connections may not be (grandma's modem connection)


- Takes server a loooong time to feed the response to grandma
- While it's doing that it can't service any more requests (进程占用资源)





Parallel Programs

- ◆ To build parallel programs, such as:
 - Parallel execution on a multiprocessor
 - Web server to handle multiple simultaneous web requests
- ◆ We will need to:
 - Create several processes that can execute in parallel
 - **Cause *each* to map to the *same* address space**
 - ▶ because they're part of the same computation
 - Give each its starting address and initial parameters
 - The OS will then schedule these processes in parallel



(b)





Processes Overheads

- ◆ A full process includes numerous things:
 - an address space (defining all the code and data pages)
 - OS resources and accounting information
 - a “thread of control”,
 - ▶ defines where the process is currently executing
 - ▶ That is the PC and registers
- ◆ Creating a new process is costly
 - all of the structures (e.g., page tables) that must be allocated
- ◆ Communicating between processes is costly
 - most communication goes through the OS





Need “Lightweight” Processes

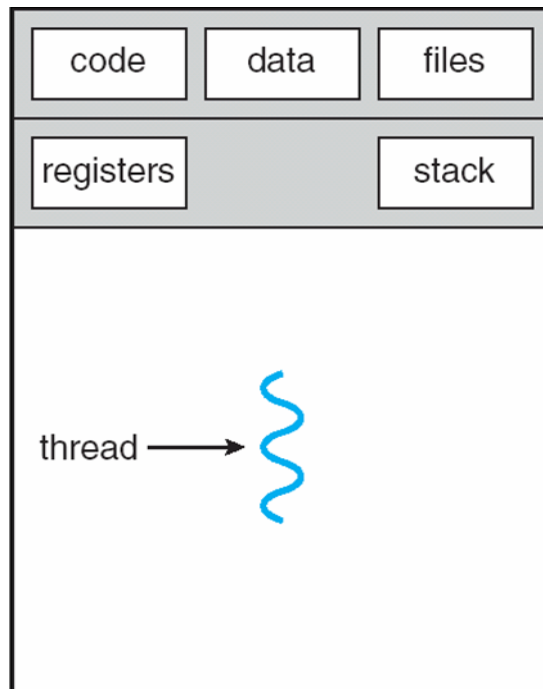
- ◆ What’s similar in these processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They share almost everything in the process
- ◆ What don’t they share?
 - Each has its own **PC, registers, and stack pointer**
- ◆ Idea: why don’t we separate the idea of process (address space, accounting, etc.) from that of the minimal “thread of control” (PC, SP, registers)?按执行线索分离






Threads and Processes

- ◆ Most operating systems therefore support two entities:
 - the process: which defines the address space and general process attributes
 - the thread: which defines a sequential execution stream within a process
- ◆ A thread is bound to a single process: For each process, however, there may be many threads.
- ◆ Threads are the unit of scheduling, Processes are *containers* in which threads execute



single-threaded process



multithreaded process





Threads vs. Processes

- ◆ A thread has no data segment or heap
 - ◆ A thread cannot live on its own, it must live within a process
 - ◆ *There can be more than one thread in a process, the first thread calls main & has the process's stack*
 - ◆ Inexpensive creation
 - ◆ Inexpensive context switching
 - ◆ If a thread dies, its stack is reclaimed
- ◆ A process has code/data/heap & other segments
 - ◆ There must be at least one thread in a process
 - ◆ Threads within a process share code/data/heap, share I/O, but each has its own stack & registers
 - ◆ Expensive creation
 - ◆ Expensive context switching
 - ◆ If a process dies, its resources are reclaimed & all threads die





Separating Threads and Processes

- ◆ Makes it easier to support multithreaded applications
 - Different from multiprocessing, multiprogramming, multitasking
- ◆ Concurrency (multithreading) is useful for:
 - improving program structure
 - handling **concurrent events** (e.g., web requests)
 - building parallel programs
 - Resource sharing
 - Multiprocessor utilization
- ◆ Is multithreading useful even on a single processor?

Benefits:

- Responsiveness
- Resource Sharing
- Economy
- Scalability





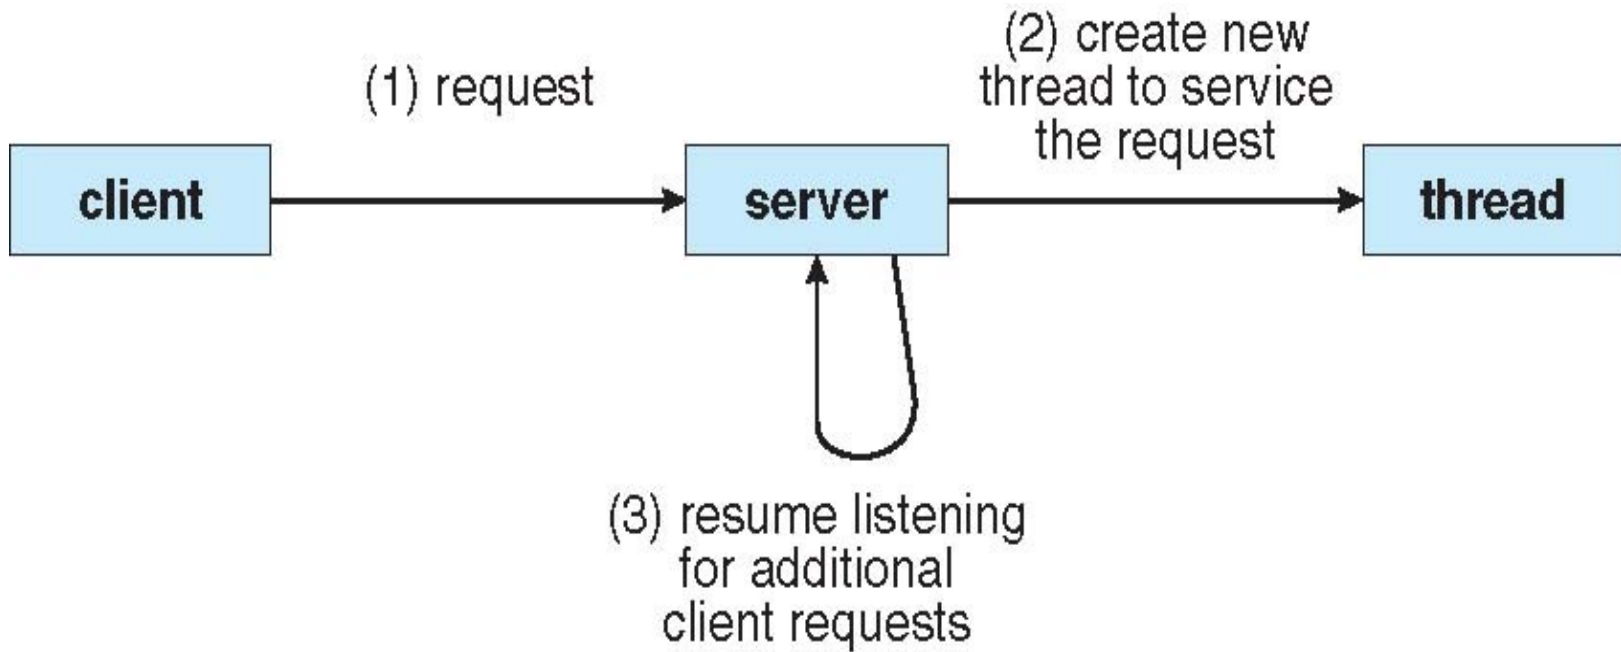
Multicore Programming

- ◆ Multicore systems putting pressure on programmers, challenges include
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**






Multithreaded Server Architecture






Concurrent Execution on a Single-core System



$$S = 1 / (1 - a + a/n)$$

Parallel Execution on a Multicore System





Amdahl's Law (阿姆达尔定律)

- ◆ S is serial portion, A is parallel portion, N processing cores

Speedup = $1 / (1 - A + A/N)$ or

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- ◆ That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- ◆ As N approaches infinity, speedup approaches $1 / S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores





Implementation of thread

There are actually 2 level of threads:

- ◆ Kernel threads:

- Supported and managed directly by the kernel.

- ◆ User threads:

- Supported above the kernel, and without kernel knowledge.





Kernel Threads

- ◆ Also called Lightweight Processes (LWP)
- ◆ Kernel threads still suffer from performance problems
- ◆ Operations on kernel threads are slow because:
 - a thread operation still requires a system call
 - kernel threads may be overly general
 - ▶ to support needs of different users, languages, etc.
 - the kernel doesn't trust the user
 - ▶ there must be lots of checking on kernel calls





User-Level Threads

- ◆ For speed, implement threads at the user level
- ◆ A user-level thread is managed by the run-time system
 - user-level code that is linked with your program
- ◆ Each thread is represented simply by:
 - PC
 - Registers
 - Stack
 - Small control block
- ◆ All thread operations are at the user-level:
 - Creating a new thread
 - switching between threads
 - synchronizing between threads





User-Level Threads

◆ User-level threads

- the thread scheduler is part of a **library**, outside the kernel
- thread context switching and scheduling is done by the library
- Can either use cooperative or pre-emptive threads
 - ▶ cooperative threads are implemented by:
 - CreateThread(), DestroyThread(), Yield(), Suspend(), etc.
 - ▶ pre-emptive threads are implemented with a timer (signal)
 - where the timer handler decides which thread to run next





Example User Thread Interface

◆ `t = thread_fork(initial context):`

create a new thread of control;

◆ `thread_stop():`

stop the calling thread, sometimes called `thread_block`;

◆ `thread_start(t):` start the named thread;

◆ `thread_yield():` voluntarily give up the processor;

◆ `thread_exit():`

terminate the calling thread, sometimes called
`thread_destroy`;





Key Data Structures

your process address space

your program:


```
for i (1, 10, I++)  
  thread_fork(I);  
....
```

your data (shared by
all your threads):


user-level thread code:

```
proc thread_fork()...  
  
proc thread_block()...  
  
proc thread_exit()...
```

queue of thread control blocks




per-thread stacks





Multiplexing User-Level Threads

- ◆ The user-level thread package sees a “virtual” processor(s)
 - it schedules user-level threads on these virtual processors
 - each “virtual” processor is implemented by a kernel thread
- ◆ The big picture
 - Create as many **kernel threads** as there are processors
 - Create as many **user-level threads** as the application needs
 - Multiplex user-level threads on top of the kernel-level threads
- ◆ Why not just create as many kernel-level threads as app needs?
 - Context switching
 - Resources





User-Level vs. Kernel Threads

User-Level

- Managed by application
- Kernel not aware of thread
- Context switching cheap
- Create as many as needed
- Must be used with care

Kernel-Level

- Managed by kernel
- Consumes kernel resources
- Context switching expensive
- Number limited by kernel resources
- Simpler to use

Key issue:

kernel threads provide virtual processors to user-level threads, but if **all of kthreads block**, then **all user-level threads will block** *even* if the program logic allows them to proceed





User/Kernel Threads

- ◆ Thread management done by user-level threads library
- ◆ Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads
- ◆ Supported by the Kernel
- ◆ Examples
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X






Multithreading Models

- ◆ Many-to-One
- ◆ One-to-One
- ◆ Many-to-Many






Many-to-One Model



Many user-level threads mapped to single kernel thread

Thread creation, scheduling, synchronization done in user space. Mainly used in language systems, portable libraries.


- ◆ Fast - no system calls required
- ◆ Few system dependencies; portable
- ◆ No parallel execution of threads - can't exploit multiple CPUs
- ◆ All threads block when one uses synchronous I/O






One-to-one Model

Each user-level thread maps to kernel thread




Thread creation, scheduling, synchronization require system calls. Used in Linux Threads, Windows NT, Windows 2000, OS/2, Solaris 9 and later

- ◆ More concurrency
- ◆ Better multiprocessor performance
- ◆ Each user thread requires creation of kernel thread
- ◆ Each thread requires kernel resources; limits number of total threads





Many-to-Many Model



- Allows many user level threads to be mapped to many kernel threads.
- Examples: Solaris, Windows NT/2000 with the *ThreadFiber* package

If $U < k$? No benefits of multithreading

If $U > k$, some threads may have to wait for an Kthread to run


- Active thread - executing on an Kthread
- Runnable thread - waiting for an Kthread

A thread gives up control of Kthread under the following:
– synchronization, lower priority, yielding, time slicing





Two-level Model



Similar to M:M, except that it allows a user thread to be **bound** to(绑定) kernel thread

- Supports both bound and unbound threads
 - Bound threads - permanently mapped to a single, dedicated kthread
 - Unbound threads - may move among kthreads in set
- Thread creation, scheduling, synchronization done in user space
- Flexible approach, "best of both worlds" (两全其美)
Used in ,Solaris 8 and earlier implementation of Pthreads and several other Unix implementations (IRIX, HP-UX, Tru64 UNIX)





Thread Libraries

- ◆ **Thread library** provides programmer with API for creating and managing threads
- ◆ Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS
- ◆ Examples:
 - Java thread
 - Win32 threads
 - POSIX Pthreads





Java Threads

- ◆ Java threads are managed by the JVM, typically implemented using the threads model provided by underlying OS:
 - On windows system, java threads are implemented using Win32 API
 - On UNIX or Linux, use Pthreads
- ◆ Java Threads包中定义了Thread类和Runnable接口, 两种创建方法:
 - ① 扩展Thread类, 并重置它的RUN()方法;
 - ② 定义一个类, 实现Runnable接口;





Extending the Thread Class

```
class Worker1 extends Thread
```

```
{ public void run() {
```

```
    System.out.println("I am a Worker Thread");
```

```
}
```

```
}
```

```
public class First
```

```
{ public static void main(String args[]) {
```

```
    Worker runner = new Worker1();
```

```
    runner.start();
```

```
    System.out.println("I am the main thread");
```

```
}
```

```
} start() 为新线程分配内存并初始化，然后调用run()方法
```






The Runnable Interface

```
public interface Runnable
{ public abstract void run();
}
class Worker2 implements Runnable
{ public void run() { /* 定义类Worker2实现Runnable接口, 定义run() 方法*/
    System.out.println("I am a Worker Thread"); }
}
public class Second
{ public static void main(String args[]) {
    Runnable runner = new Worker2();
    Thread thrd = new Thread(runner);
    thrd.start();
    System.out.println("I am the main thread");
} /* 创建线程对象传递给Runnable对象, 新线程由start() 方法创建后,
    开始执行Runnable对象的run() 方法*/
```





Java Thread States





Win32 Threads

- Win32 API is the primary API for Microsoft OS (Win95,98,NT,2000,XP),A kernel-level library on windows systems

Windows XP Threads:

- Implements the **one-to-one mapping, kernel-level**
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The register set, stacks, and private storage area are known as the **context** of the threads

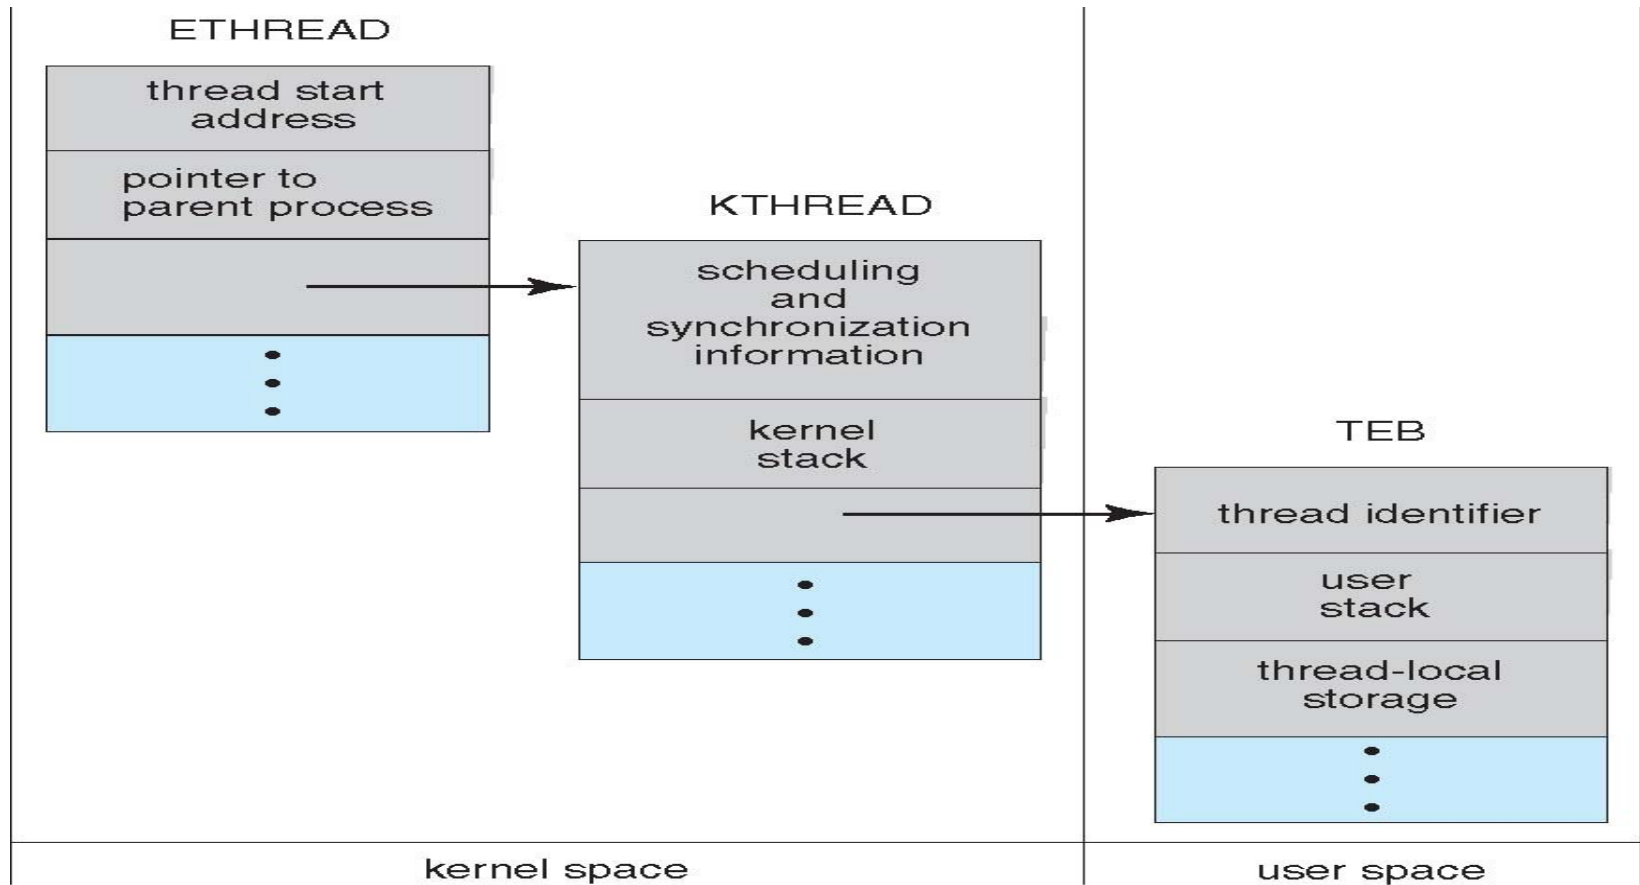




Windows XP Threads

■ The primary data structures of a thread include:

- ETHREAD (executive thread block)
- KTHREAD (kernel thread block)
- TEB (thread environment block)





Process Control Block (PCB)

Process ID (PID)
Parent PID
...
Next Process Block •
List of open files •
Image File Name
List of Thread Control Blocks •
...


- This is an abstract view
- Windows implementation of PCB is split in multiple data structures

PCB

Handle Table

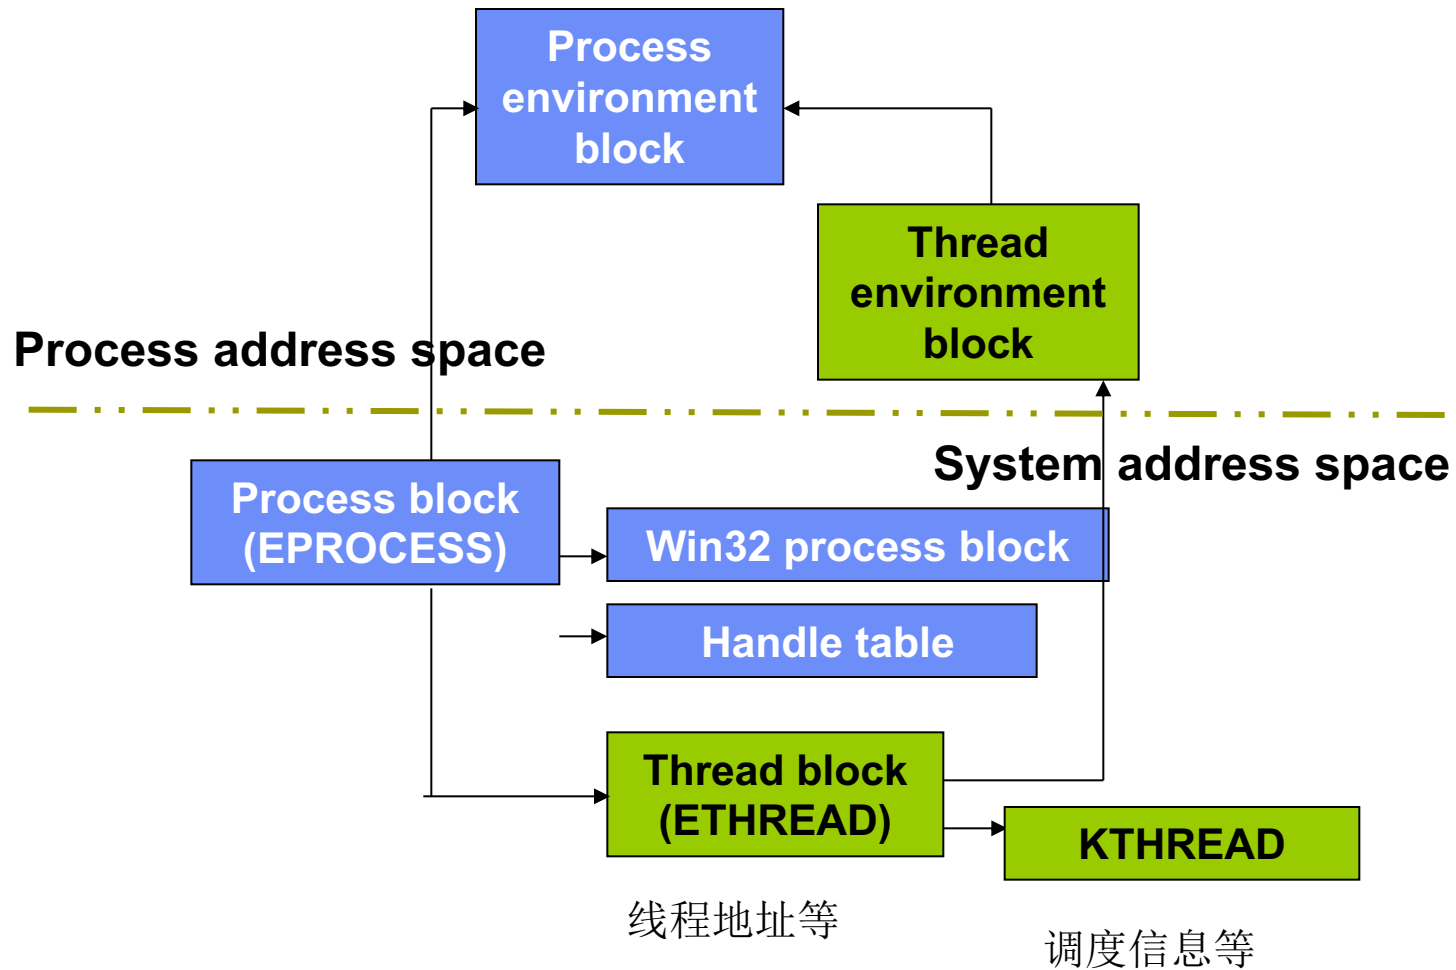
Thread Control Block (TCB)

Next TCB •
Program Counter
Registers
...





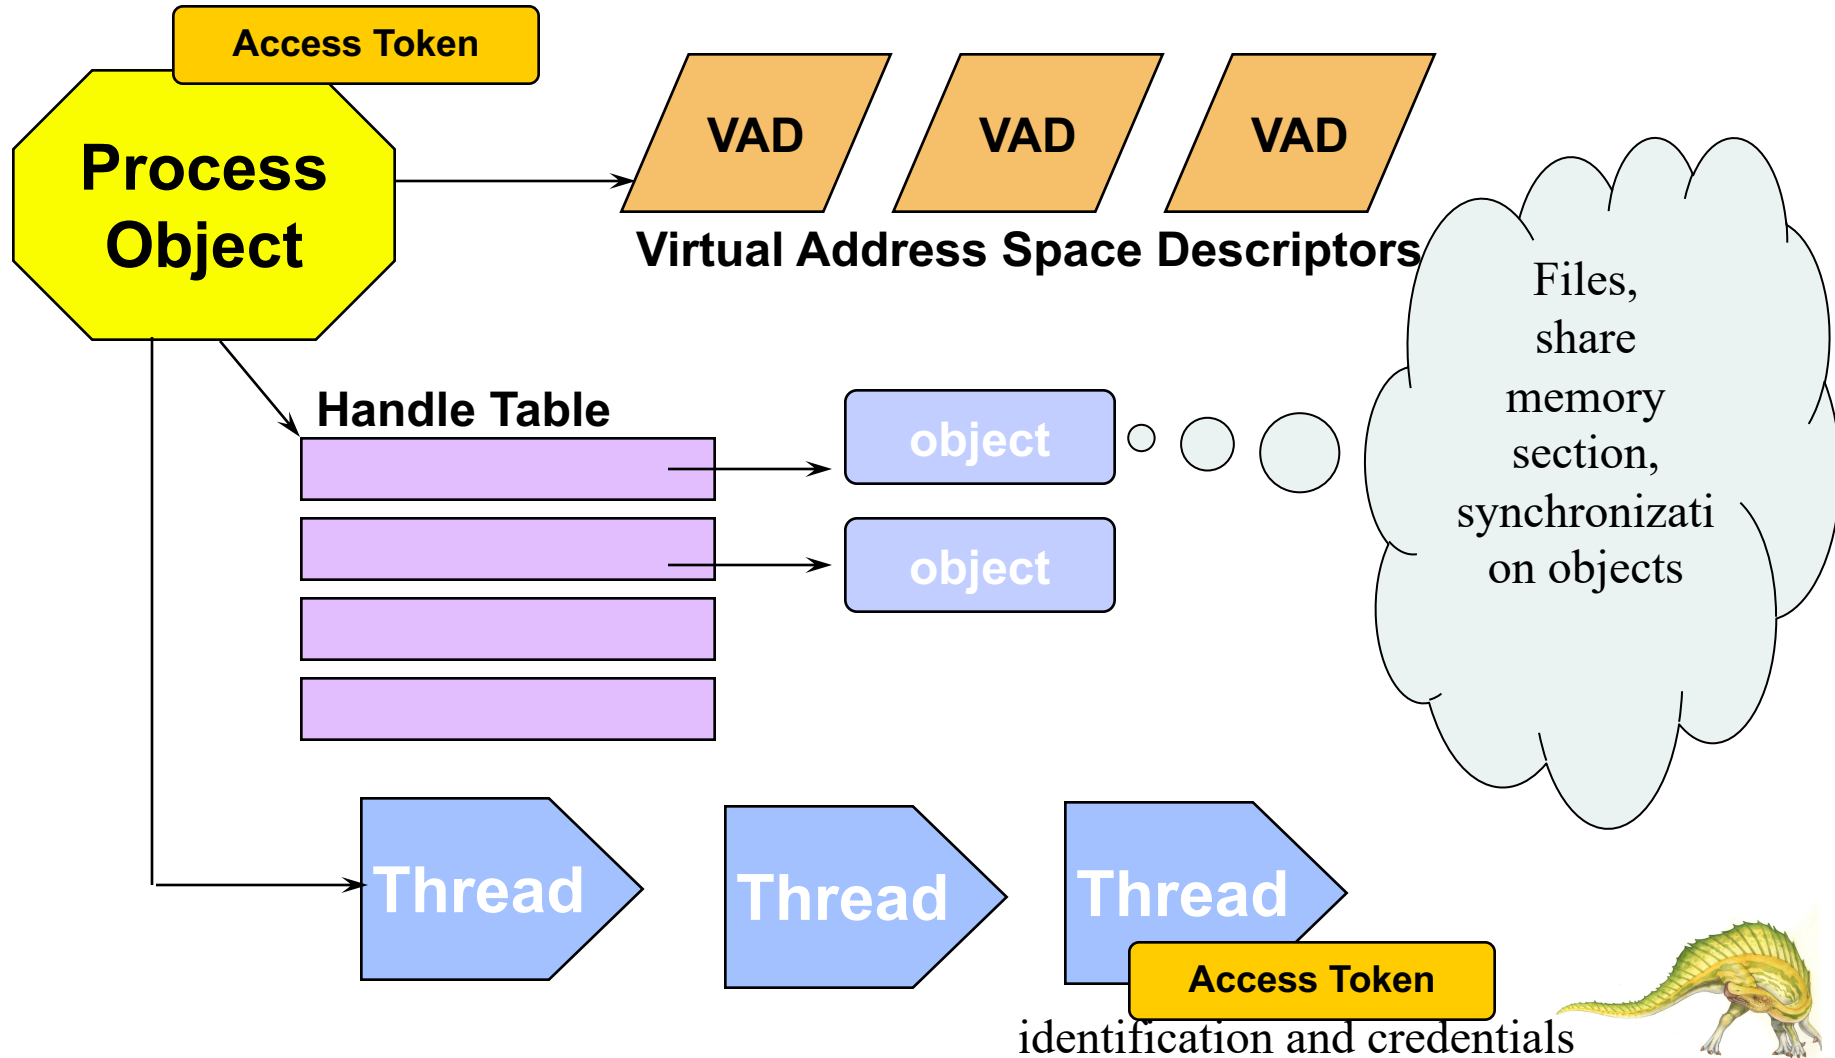
Windows Process and Thread Internals





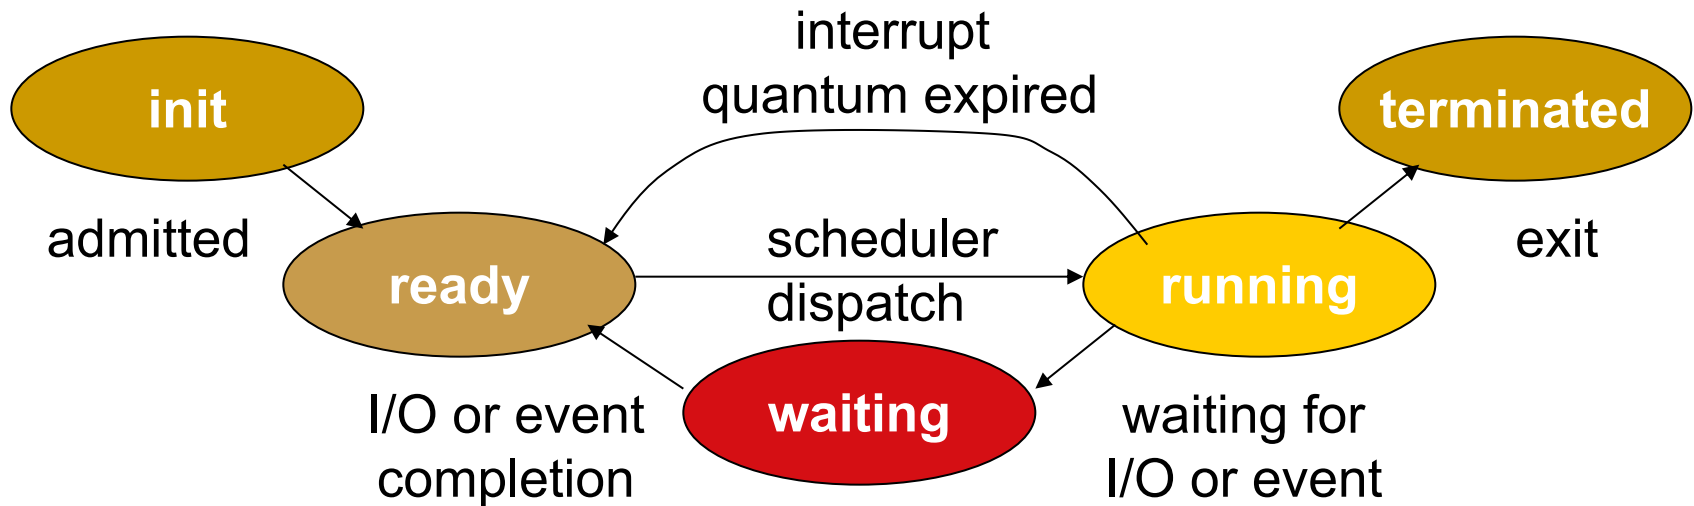
Windows Processes & Threads

Internal Data Structures



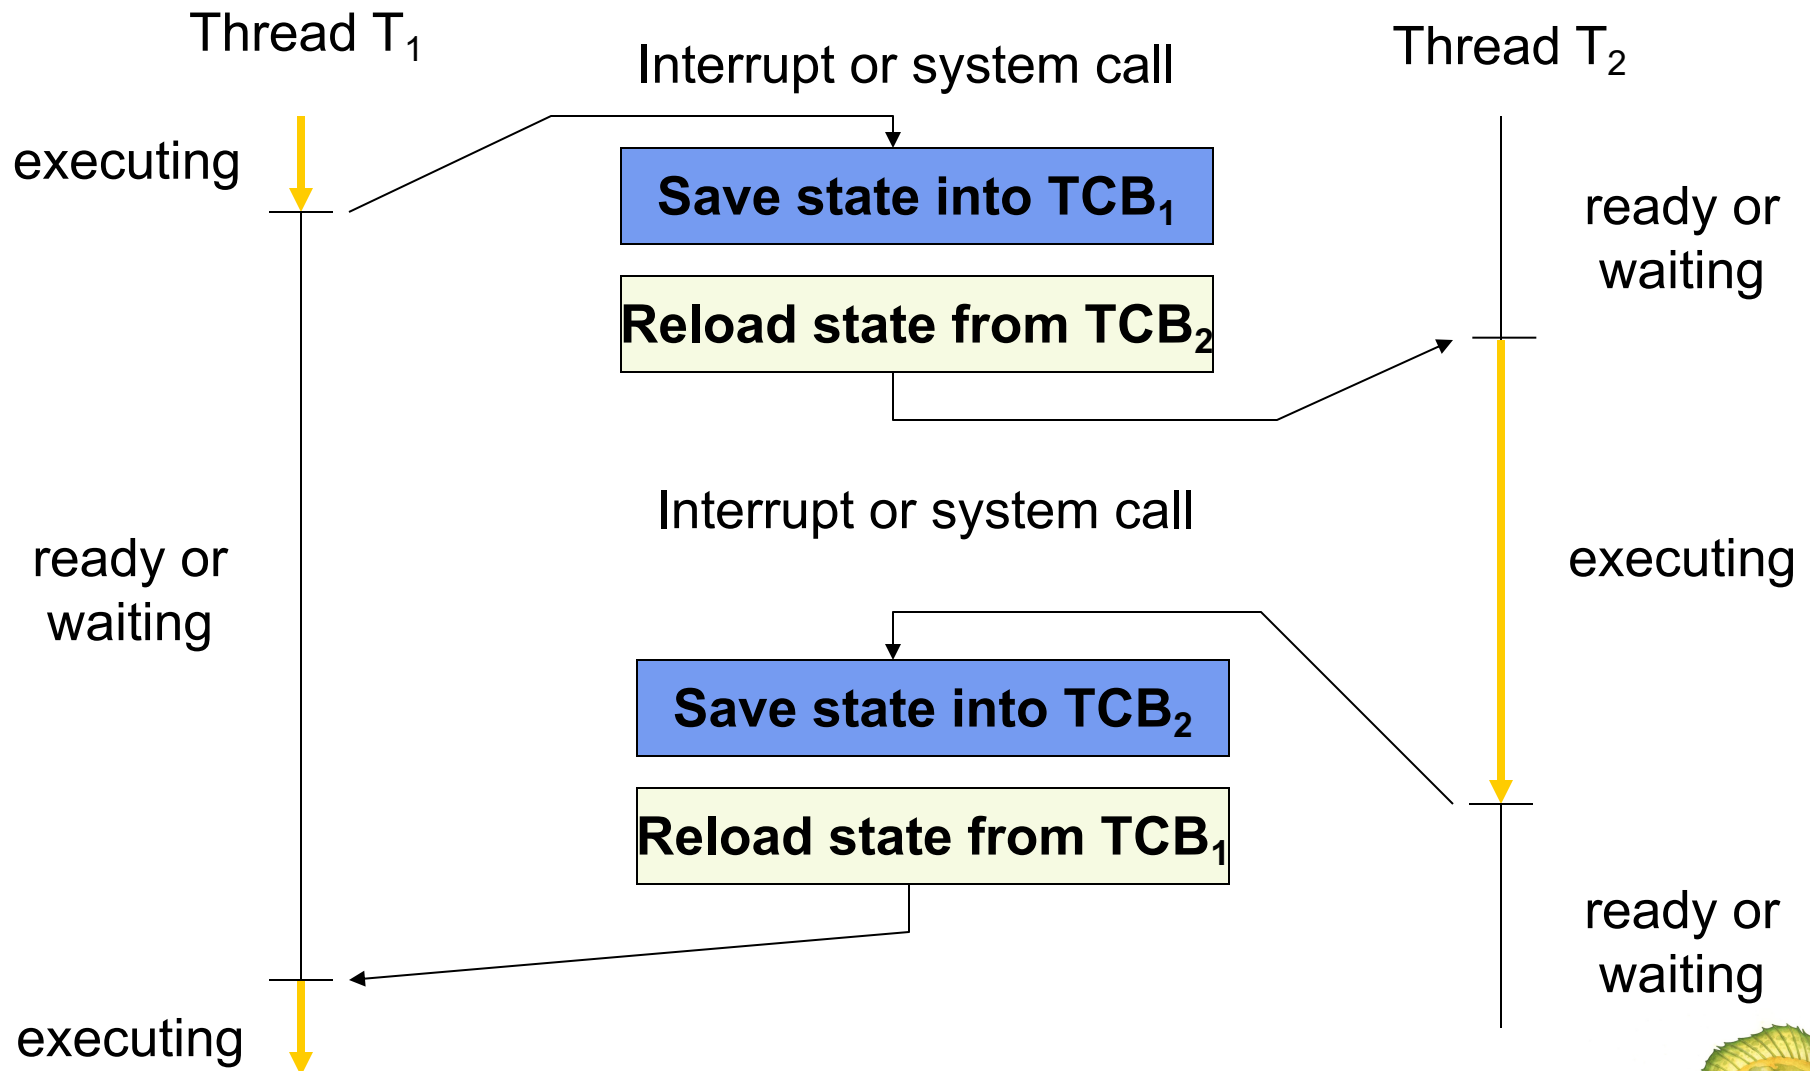


Windows XP Threads





CPU Switch from Thread to Thread





POSIX Pthreads

- ◆ May be provided either as user-level or kernel-level
- ◆ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ◆ API specifies behavior of the thread library, implementation is up to development of the library
- ◆ Common in UNIX operating systems
(Linux, Solaris, Mac OS X)





Linux Threads

- ◆ Linux does not distinguish between processes and threads; Linux refers to them as *tasks* rather than *process* or *thread*;
- ◆ Thread creation is done through system call `clone()` or `pthread_create()` ;
- ◆ 调用 `fork()` 创建一个新进程, 它具有父进程所有相关数据结构的拷贝;
- ◆ 调用 `clone()` 创建新进程, 但它不复制父进程的数据结构, 而是指向了父进程的数据结构, 从而允许子进程共享父进程的内存和其他资源。或者定义一个函数作为线程的入口, 然后调用 `pthread_create()` 创建线程。





Linux Threads

- ◆ **clone()** allows a child task to share the address space of the parent task (process)
- ◆ **Flags** determine how much sharing is to take place between the parent and child
- ◆ If no flag is set when clone(), no sharing take place, resulting in functionality similar to fork()


flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.





Solaris threads

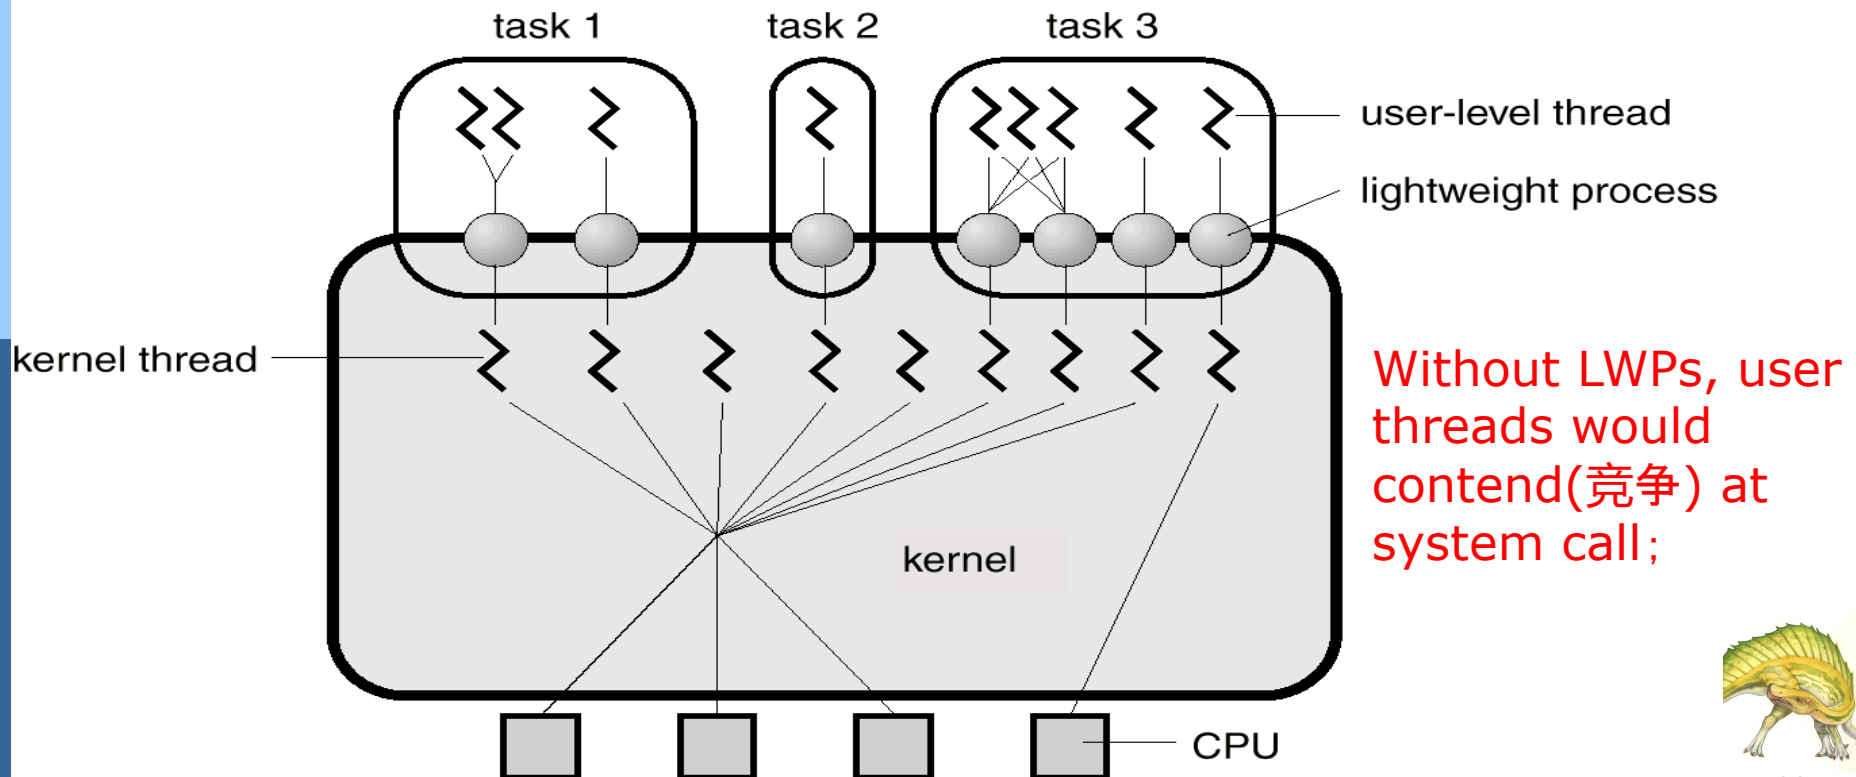
- ◆ Solaris is a version of UNIX with support for threads at the kernel and user levels, symmetric multiprocessing, and real-time scheduling.
 - Kernel threads;
 - Lightweight Processes;
 - User Level Threads;
- ◆ LWP – (lightweight processes) intermediate level between user-level threads and kernel-level threads.





Solaris threads

- ◆ **Kernel Threads:** small data structure and a stack; thread switching relatively fast. Kernel only sees the LWPs that support user-level threads.
- ◆ **LWP:** a register set for the user-thread it is running, accounting and memory information, switching between LWPs is relatively slow.
- ◆ **User Threads:** only need stack and program counter; no kernel involvement means fast switching.

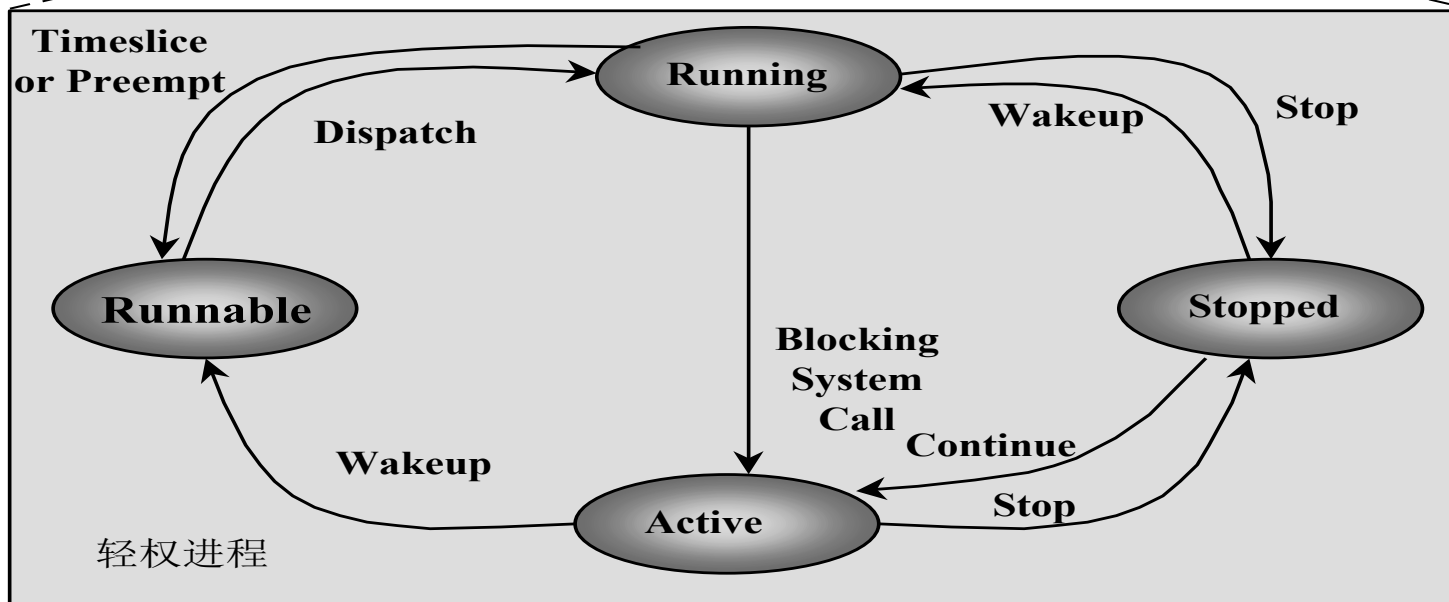
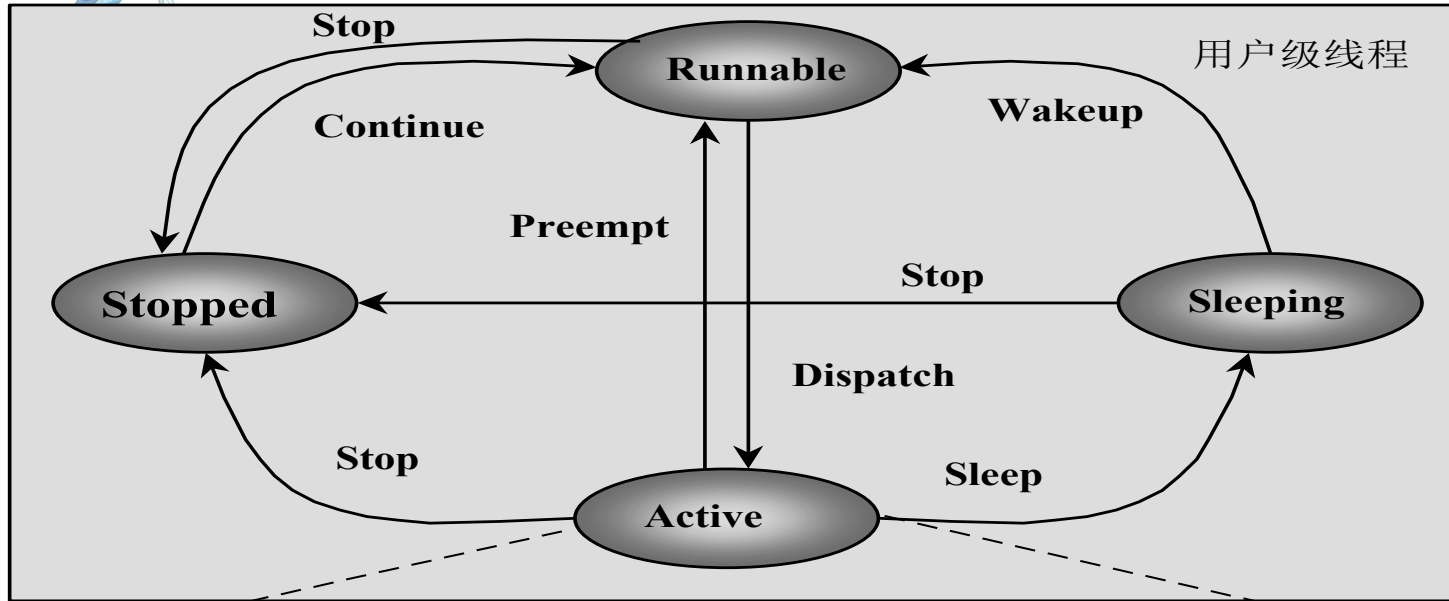




Solaris threads

- ◆ User level threads may be either bound or unbound;
- ◆ Bound: A user thread is permanently attached to a LWP;
- ◆ All unbound threads in an application are multiplexed onto the pool of available for the application;
 - Threads are unbound by default .
- ◆ The thread library adjusts LWPs in the pool
 - The thread library ages LWPs and deletes them when they are unused for a long time, typically about 5 minutes.





Solaris: User Threads and LWP





C library functions on Solaris threads

•Create a User Threads

- **int thr_create**(void *stack_base, size_t stack_size, void *(*start_routine)(void *), void *arg, long flags, thread_t *new_thread_id);

flags: THR_BOUND(永久捆绑)

THR_NEW_LWP(创建新LWP放入LWP池)

两者同时指定则创建两个新LWP，一个永久捆绑而另一个放入LWP池；

•Create a LWP

int _lwp_create(ucontext_t *contextp, unsigned long flags, lwpid_t *new_lwp_id);

• Make a context on LWP

- **void _lwp_makecontext**(ucontext_t *ucp, void (*start_routine)(void *), void *arg, void *private, caddr_t stack_base, size_t stack_size);





Implicit Threading

- ◆ Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- ◆ Creation and management of threads done by compilers and run-time libraries rather than programmers
- ◆ Three methods explored
 - Thread Pools
 - OpenMP
 - Grand Central Dispatch
- ◆ Other methods include Microsoft Threading Building Blocks (TBB), **java.util.concurrent** package





Thread Pools

- ◆ Create a number of threads in a pool where they await work
- ◆ Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
 - Separating task to be performed from mechanics of creating task allows different strategies for running task (执行与创建分离)
 - ▶ i.e. Tasks could be scheduled to run periodically
- ◆ Windows API supports thread pools:

```
DWORD WINAPI PoolFunction(AVOID Param) {  
    /*  
     * this function runs as a separate thread.  
     */  
}
```





OpenMP

- ◆ Set of compiler directives and an API for C, C++, FORTRAN
- ◆ Provides support for parallel programming in shared-memory environments
- ◆ Identifies **parallel regions** – blocks of code that can run in parallel

#pragma omp parallel

Create as many threads as there are cores

#pragma omp parallel for

for(i=0;i<N;i++) {

c[i] = a[i] + b[i];

}

Run for loop in parallel

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```





Grand Central Dispatch

- ◆ Apple technology for Mac OS X and iOS operating systems
- ◆ Extensions to C, C++ languages, API, and run-time library
- ◆ Allows identification of parallel sections
- ◆ Manages most of the details of threading
- ◆ Block is in “`^ { }`” - `^ { printf("I am a block"); }`
- ◆ Blocks placed in dispatch queue
 - Assigned to available thread in thread pool when removed from queue





Grand Central Dispatch

◆ Two types of dispatch queues:

- serial – blocks removed in FIFO order, queue is per process, called **main queue**
 - ▶ Programmers can create additional serial queues within program
- concurrent – removed in FIFO order but several may be removed at a time
 - ▶ Three system wide queues with priorities low, default, high

```
dispatch_queue_t queue = dispatch_get_global_queue  
    (DISPATCH_QUEUE_PRIORITY_DEFAULT, 0);
```

```
dispatch_async(queue, ^{ printf("I am a block."); });
```





Threading Issues

- ◆ Semantics of **fork()** and **exec()** system calls
- ◆ Thread cancellation of target thread
 - Asynchronous or deferred
- ◆ Signal handling
- ◆ Thread pools
- ◆ Thread-specific data
- ◆ Scheduler activations





Thread Hazards

```
int a = 1, b = 2, w = 2;
```

```
main() {  
    CreateThread(fn, 4);  
    CreateThread(fn, 4);  
    while(w) ;  
}
```

```
fn() { int v = a + b;  
      w--;  
}
```

- A statement like “w--” in C (or C++) is implemented by several machine instructions:

```
ld      r4, #w  
add     r4, r4, -1  
st      r4, #w
```

- Now, imagine the following sequence, what is the value of w?

```
ld      r4, #w  
_____  
_____  
_____  
add     r4, r4, -1  
st      r4, #w
```

```
ld      r4, #w  
add     r4, r4, -1  
st      r4, #w
```





作业

- 编写一个C/S架构的分布式程序, Server接收Client发来的请求, 执行一个计算 $F(X)$ 并给Client返回结果; 分别用进程与线程作为服务器Server实现, 并比较服务器的开销. 可以在一台机器上模拟.



End of Chapter 4

