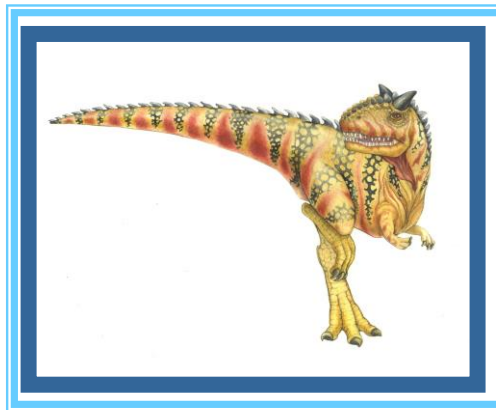


# Chapter 7: Deadlocks

---





# Chapter 7: Deadlocks

---

- ◆ The Deadlock Problem
- ◆ System Model
- ◆ Deadlock Characterization
- ◆ Methods for Handling Deadlocks
- ◆ Deadlock Prevention
- ◆ Deadlock Avoidance
- ◆ Deadlock Detection
- ◆ Recovery from Deadlock





# Chapter Objectives

---

- ◆ To develop a description of deadlocks, which prevent sets of concurrent processes from completing their tasks;
- ◆ To present a number of different methods for preventing or avoiding deadlocks in a computer system;





# The Deadlock Problem



- ◆ A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- ◆ The resources may be either physical resources or logical resources.





# The Deadlock Problem

## ◆ Example

- ✓ System has 2 disk drives
- ✓  $P_1$  and  $P_2$  each hold one disk drive and each needs another one

## ◆ Example

- ✓ semaphores  $A$  and  $B$ , initialized to 1


$P_0$   
wait ( $A$ );  
wait ( $B$ );

$P_1$   
wait( $B$ )  
wait( $A$ )





# Deadlock Example




(a) A potential deadlock

(b) an actual deadlock





# Bridge Crossing Example



- ◆ Traffic only in one direction
- ◆ Each section of a bridge can be viewed as a resource
- ◆ If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
- ◆ Several cars may have to be backed up if a deadlock occurs
- ◆ Starvation is possible
- ◆ Note – Most OSes do not prevent or deal with deadlocks





# Why Deadlock?

---

- ◆ Resource competition
  - ✓ The number of resources is not sufficient for processes requirements
- ◆ Advance order of Concurrent processes is not appropriate. In a multiprogramming system, it is asynchronous for concurrent processes execution
  - ✓ In some order, each process can complete successfully
  - ✓ In some order, there will be a deadlock








**Concurrent processes P1,P2; resources R1 and R2 has only one instance respectively**

**Curve4 will lead to unsafe area**






# System Model

---

- ◆ Resource types  $R_1, R_2, \dots, R_m$ 
  - ◆ *CPU cycles, memory space, I/O devices*
- ◆ Each resource type  $R_i$  has  $W_i$  instances.
- ◆ The number of resources requested may not exceed the total number of resources available in the system.





# How process utilizes a resource?

---

- ◆ Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **Release**
- ◆ The request and release are system calls, such as `open()` and `close()` file, and `allocate()` and `free()` memory system calls.





# Deadlock with Semaphores

## ◆ Data:

- A semaphore  $S_1$  initialized to 1
- A semaphore  $S_2$  initialized to 1

## ◆ Two threads $T_1$ and $T_2$

### ◆ $T_1$ :

`wait(s1)`

`wait(s2)`

### ◆ $T_2$ :

`wait(s2)`

`wait(s1)`





# Necessary Conditions for Deadlock

Deadlock can arise if four conditions hold **simultaneously**.


- ◆ **Mutual exclusion:** only one process at a time can use a resource
- ◆ **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- ◆ **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task





# Necessary Conditions for Deadlock(Cont)

- ◆ **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2, \dots, P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .





# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

◆  $V$  is partitioned into two types:

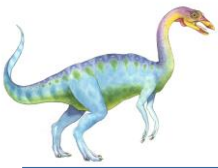
✓  $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system

✓  $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system

◆ request edge – directed edge  $P_1 \rightarrow R_j$

◆ assignment edge – directed edge  $R_j \rightarrow P_i$







# Resource-Allocation Graph



(a)



(b)



(c)

- ➡ Resource R assigned to process A
- ➡ Process B is requesting/waiting for resource S
- ➡ Process C and D are in deadlock over resources T and U







# Resource-Allocation Graph (Cont.)


◆ Process




◆ Resource Type with 4 instances



◆  $P_i$  requests instance of  $R_j$




◆  $P_i$  is holding an instance of  $R_j$






# Resource Allocation Graph


## Example of a Resource Allocation Graph



## Resource Allocation Graph With A Deadlock



## Graph With A Cycle But No Deadlock





# Basic Facts

---

- ◆ If graph contains no cycles  $\Rightarrow$  no deadlock
- ◆ If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock





# Methods for Handling Deadlocks

---

- ◆ Ensure that the system will *never* enter a deadlock state
  - Deadlock prevention (静态预防)
  - Deadlock avoidance (动态避免)
- ◆ Allow the system to enter a deadlock state and then recover
  - Deadlock detection and recovery (动态消除)
- ◆ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX: **Ostrich policy**





# Deadlock Prevention

---

◆ Restrain the ways request can be made, so that at least one of **the necessary conditions for deadlock can not occur.**

◆ Possible side effects:

- Low device utilization
- Reduced system throughput

◆ **Mutual Exclusion**

- not required for sharable resources; must hold for nonsharable resources.





# Deadlock Prevention

---


- ◆ **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution
  - or allow process to request resources only when the process has none
- ◆ Low resource utilization; starvation possible





# Dining Philosophers

- ◆ 5 philosophers sitting around a round table
- ◆ 1 chopstick in between each pair of philosophers
  - 5 chopsticks total
- ◆ Each philosopher needs two chopsticks to eat
- ◆ How to prevent deadlock





# Deadlock Prevention (Cont.)

## ◆ No Preemption –

### 方案1

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

### 方案2

- When process A requests some unavailable resources, we check whether they are allocated to another process B that is waiting for other resources;
- If so, we preempt the desired resources from B and allocate them to A.
- Otherwise, A must wait.








# Deadlock Prevention (Cont.)

- ◆ Allow Preemption!
- ◆ Can preempt CPU by
  - saving its state to thread control block and resuming later
- ◆ Can preempt memory by
  - swapping memory out to disk and loading it back later
- ◆ Can we preempt the holding of a lock?
- ◆ Some resource cannot be preempted
- ◆ Consider a process given the printer
  - halfway through its job
  - now forcibly take away printer
  - !!??






# Deadlock Prevention (Cont.)

## ◆ Circular Wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

For example, Dining Philosophers





# Deadlock Avoidance

---

Requires that the system has some additional *a priori* information available. E.g, which type resources a process needs? How many instances of each resource type a process needs at most?





# Deadlock Avoidance

- ◆ Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need;
- ◆ The deadlock-avoidance algorithm dynamically examines the resource-allocation *state* to ensure that there can *never be a circular-wait condition*;
- ◆ Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes;
- ◆ When a process requests an available resource, system must decide if immediate allocation leaves the system in a *safe state*.





# Safe State

- ◆ System is in safe state if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$ , for each  $P_i$  the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ .
- ◆ That is:
$$\langle P_1, P_2, \dots, P_j, P_i, \dots, P_n \rangle$$
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on





## example

◆ processes:  $P_1$ ,  $P_2$  and  $P_3$ . 12 tapes. Snapshot at  $T_0$  :

◆ process	max	allocated	need	available
◆ $P_1$	10	5	5	3
◆ $P_2$	4	2	2	
◆ $P_3$	9	2	7	

◆ sequence ( $P_2$ ,  $P_1$ ,  $P_3$ ) satisfies the safety condition :

- Allocate 2 of the 3 available tapes to  $P_2$ , which satisfies the max need of  $P_2$ . when  $P_2$  finishes, the system will have 5 available tapes.
- Then allocate the 5 tapes to  $P_1$ . when  $P_1$  finishes, the system will have 10 available tapes.
- At last, allocate 7 of the 10 tapes to  $P_3$
- **So, system is safe at  $T_0$ .**





## System goes from a safe state to an unsafe state

- ◆ Suppose that, at  $T_0$ ,  $P_3$  requests a tape and is allocated one tape. Then system goes from  $T_0$  to  $T_1$ .

Process	max	allocated	need	available	available
				before allocating	after releasing
$P_1$	10	5	5	>	
$P_2$	4	2	2	=<	2
$P_3$	9	3	6	>	4

- ◆ At  $T_1$ , we can not find a sequence satisfying safety condition, so, it is unsafe.
- ◆ Therefore, when  $P_3$  requests a tape at  $T_0$ , we must make it wait to avoid deadlock.






# Basic Facts

- ◆ If a system is in safe state  $\Rightarrow$  no deadlocks
- ◆ If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- ◆ Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

**Safe, Unsafe ,  
Deadlock State**







# Avoidance algorithms

---

- ◆ Single instance of a resource type
  - Use a resource-allocation graph
  
- ◆ Multiple instances of a resource type
  - Use the banker's algorithm






# Resource-Allocation Graph Scheme

- ◆ **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_j$  **may request** resource  $R_j$ ; represented by a dashed line
- ◆ Claim edge converts **to request edge** when a process requests a resource
- ◆ Request edge converted to an **assignment edge** when the resource is allocated to the process
- ◆ When a resource is released by a process, assignment edge reconverts to a **claim edge**
- ◆ Resources must be claimed *a priori* in the system







## Resource-Allocation Graph



## Unsafe State In Resource-Allocation Graph





# Resource-Allocation Graph Algorithm

---

- ◆ Suppose that process  $P_i$  requests a resource  $R_j$
- ◆ The request can be granted only if **converting the request edge to an assignment edge** does not result in the formation of a **cycle** in the resource allocation graph





# Banker's Algorithm

---

- ◆ Multiple instances
- ◆ Each process must a priori claim maximum use
- ◆ When a process requests a resource it may have to wait
- ◆ When a process gets all its resources it must return them in a finite amount of time





# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes,  $m$  = number of resources types.

- ◆ **Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- ◆ **Max:**  $n \times m$  matrix. If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- ◆ **Allocation:**  $n \times m$  matrix. If  $Allocation[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- ◆ **Need:**  $n \times m$  matrix. If  $Need[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$





# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

*Work* = *Available*

*Finish* [*i*] = *false* for *i* = 0, 1, ..., *n* - 1

2. Find an *i* such that both:

(a) *Finish* [*i*] = *false*

(b)  $Need_i \leq Work$

If no such *i* exists, go to step 4

3. *Work* = *Work* + *Allocation*<sub>*i*</sub>  
*Finish*[*i*] = *true*  
go to step 2

4. If *Finish* [*i*] == *true* for all *i*, then the system is in a safe state





# Resource-Request Algorithm for Process $P_i$

$Request$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$Available = Available - Request_i;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored







# Example of Banker's Algorithm

- ◆ 5 processes  $P_0$  through  $P_4$ ;
- 3 resource types: A (10 instances)
- B (5 instances), and C (7 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

- ◆  $Need =$   
 $Max - Allocation$

	<u>Need</u>
	A B C
$P_0$	7 4 3
$P_1$	1 2 2
$P_2$	6 0 0
$P_3$	0 1 1
$P_4$	4 3 1





## Example (Cont.)

- ◆ To find a safe sequence using safety algorithm

Work[]=available=(3,3,2); Finish[i]=false (i=0..4);

	Work	need	allocation	finish
P1	3 3 2	1 2 2	2 0 0	T
P3	5 3 2	0 1 1	2 1 1	T
P4	7 4 3	4 3 1	0 0 2	T
P2	7 4 5	6 0 0	3 0 2	T
P0	10 4 7	7 4 3	0 1 0	T

- ◆ The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria





## Example (Cont.)

- ◆ Usually, there would be more than one safe sequences if the system is in safe state.

Work[]=available=(3,3,2); Finish[i]=false (i=0..4)

	Work	need	allocation	finish
P3	3 3 2	0 1 1	2 1 1	T
P4	5 4 3	4 3 1	0 0 2	T
P1	5 4 5	1 2 2	2 0 0	T
P2	7 4 5	6 0 0	3 0 2	T
P0	10 4 7	7 4 3	0 1 0	T

- Safe sequence: (P3, P4, P1, P2, P0)





## Example: $P_1$ Request (1,0,2)

- ◆ Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true)

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	1	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- ◆ Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- ◆ Can request for (3,3,0) by  $P_4$  be granted?
- ◆ Can request for (0,2,0) by  $P_0$  be granted?





# Deadlock Detection

---

- ◆ Allow system to enter deadlock state
- ◆ Detection algorithm
- ◆ Recovery scheme






# Single Instance of Each Resource Type

- ◆ Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- ◆ Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- ◆ An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph






# Resource-Allocation Graph and Wait-for Graph




(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph





# Several Instances of a Resource Type

---

- ◆ **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- ◆ **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- ◆ **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i_j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .







# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively Initialize:
  - (a) *Work* = *Available*
  - (b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
 $Finish[i] = false$ ; otherwise,  $Finish[i] = true$
2. Find an index  $i$  such that both:
  - (a)  $Finish[i] == false$
  - (b)  $Request_i \leq Work$If no such  $i$  exists, go to step 4
3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked

**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state**





# Example of Detection Algorithm

- ◆ Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- ◆ Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- ◆ Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$





## Example (Cont.)

- ◆  $P_2$  requests an additional instance of type C

### Request

	A	B	C
$P_0$	0	0	0
$P_1$	2	0	1
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- ◆ State of system?

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes requests
- Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$





# Detection-Algorithm Usage

---

- ◆ When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - one for each disjoint cycle

Deadlock occurs only when some process makes a request that can not be granted immediately:

- ◆ We can invoke deadlock detection algorithm every time a resource request cannot be granted immediately.
  - Lead to a considerable overhead in computation time.
- ◆ A less expensive alternative is to invoke the algorithm at less frequent intervals.
  - Once per hour
  - Whenever CPU utilization drops below 40%.





# Recovery from Deadlock: Process Termination

---

- ◆ Abort all deadlocked processes
- ◆ Abort one process at a time until the deadlock cycle is eliminated
- ◆ In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?





# Recovery from Deadlock: Resource Preemption

---

- ◆ Selecting a victim – minimize cost
- ◆ Rollback – return to some safe state, restart process for that state
- ◆ Starvation – same process may always be picked as victim, include number of rollback in cost factor





# assignment

---

1. 7.1
2. 7.9
3. 7.13

