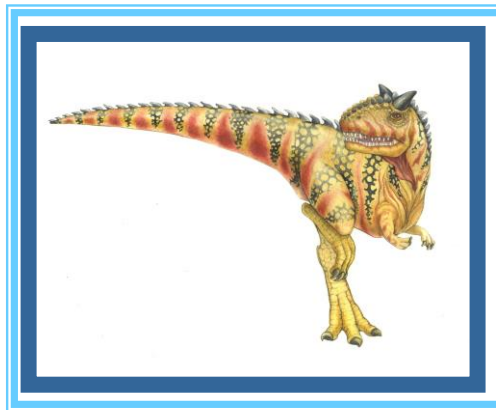


Chapter 9: Virtual Memory





Chapter 9: Virtual Memory

- ◆ Background
- ◆ Demand Paging
- ◆ Copy-on-Write (写时复制)
- ◆ Page Replacement
- ◆ Allocation of Frames
- ◆ Thrashing
- ◆ Memory-Mapped Files
- ◆ Allocating Kernel Memory
- ◆ Other Considerations
- ◆ Operating-System Examples





Objectives

- ◆ To describe the benefits of a virtual memory system
- ◆ To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- ◆ To discuss the principle of the working-set model





Background

- ◆ **Virtual memory** is a technique that allows the execution of processes that may not be completely in memory. Separation of user logical memory from physical memory.
 - ▣ Only part of the program needs to be in memory for execution
 - ▣ Logical address space can therefore be much larger than physical address space
 - ▣ Allows address spaces to be shared by several processes
 - ▣ Allows for more efficient process creation





Background

◆ principle of locality:

- **Temporal Locality** (Locality in Time):同一事物的访问在时间上聚集在一起;
- **Spatial Locality** (Locality in Space):在时间上被引用的事物在空间上也接近(相邻的存储器地址, 磁盘上的邻近扇区等), 则引用序列被认为具有**空间局部性 (Spatial Locality)**。

- ✓ sum为时间局限性
- ✓ 数组元素为空间局限性

```
// Compute sum of an int array
int  a[N] = {2, 5, 3, 7, ...};
int  sum=0;

for(i=0; i<N; i++)
    sum = sum + a[i];
```


◆ Virtual memory can be implemented via:

- Demand paging
- Demand segmentation





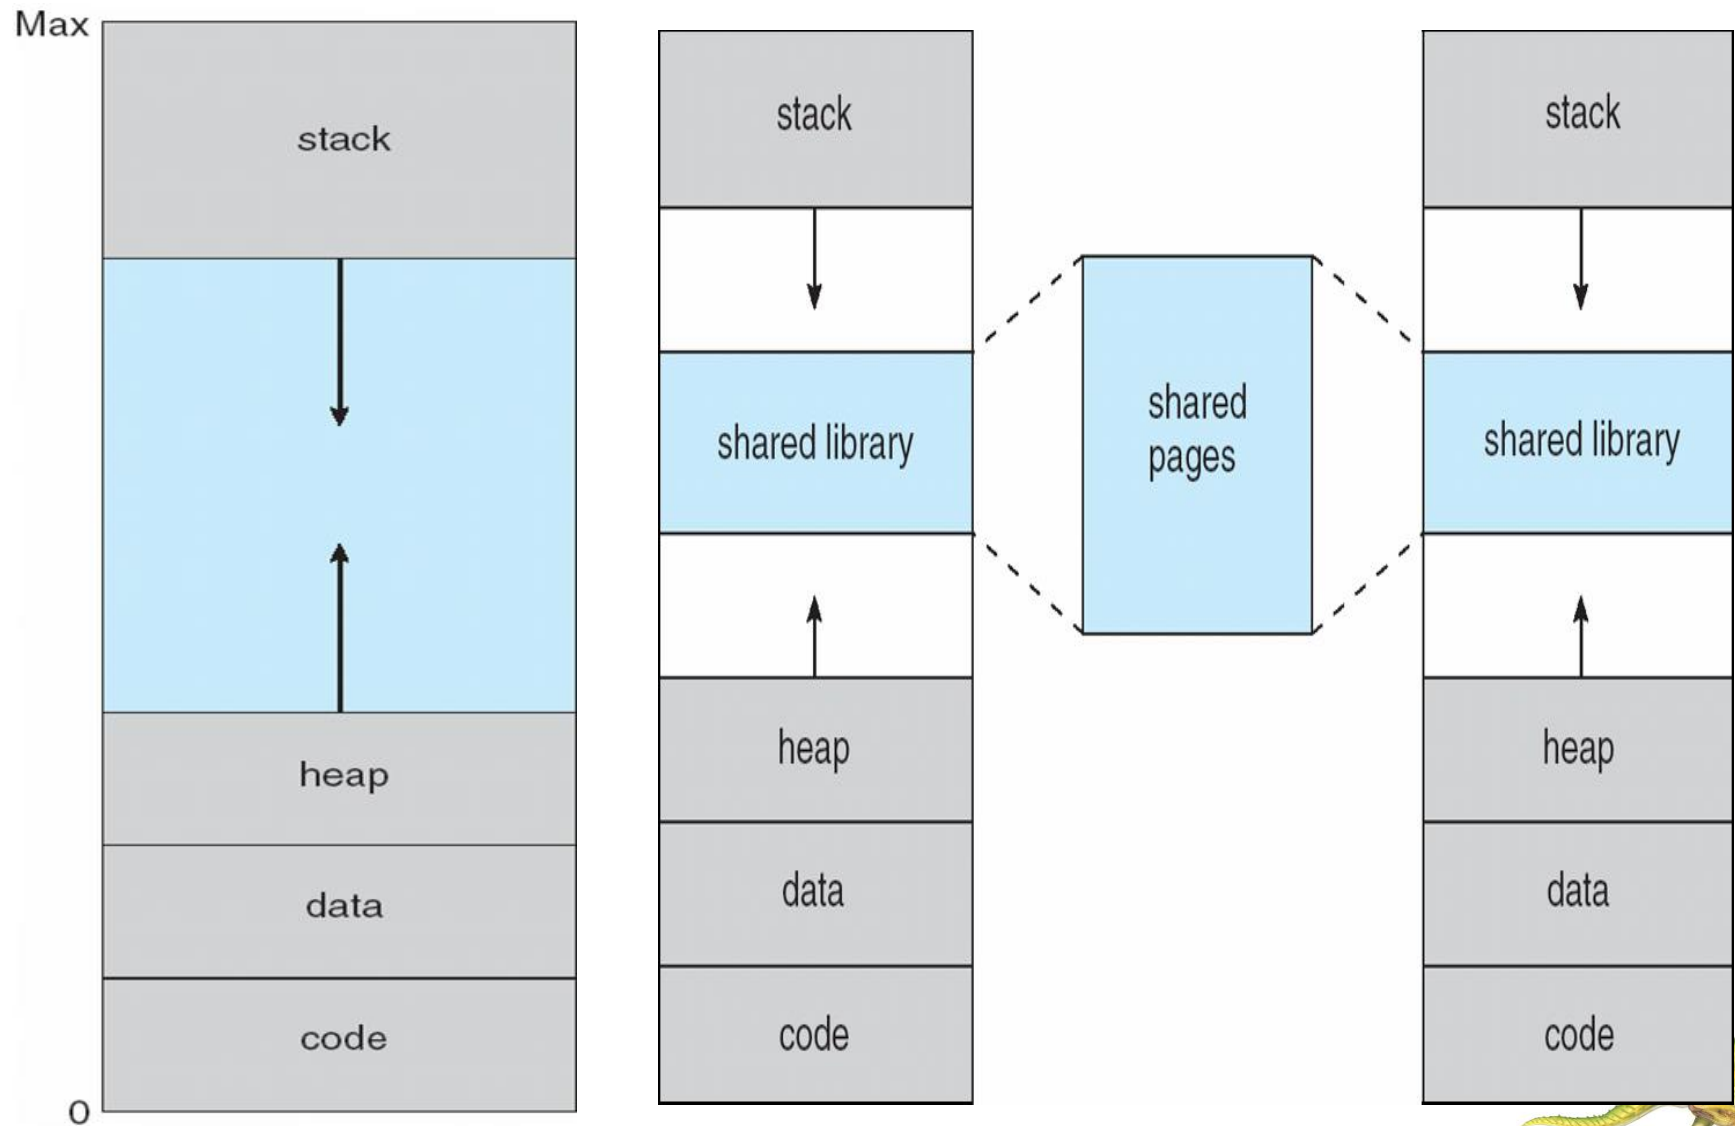
Virtual Memory That is Larger Than Physical Memory





Virtual-address Space

Shared Library Using Virtual Memory






Demand Paging

- ◆ Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- ◆ Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- ◆ **Lazy swapper** – never swaps a page into memory unless page will be needed
 - Swapper that deals with pages is a **pager**





Transfer of a Paged Memory to Contiguous Disk Space





Valid-Invalid Bit

- ◆ With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory, **i** \Rightarrow not-in-memory)
- ◆ Initially valid–invalid bit is set to **i** on all entries
- ◆ Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

- ◆ During address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault





Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory

valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory



Page 0,2,5 in-memory





Page Fault

- ◆ If there is a reference to a page, first reference to that page will trap to operating system:

page fault

- ◆ Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
- ◆ Get empty frame
- ◆ Swap page into frame
- ◆ Reset tables
- ◆ Set validation bit = **v**
- ◆ Restart the instruction that caused the page fault





Page Fault

缺页中断处理代码示例：

// 定义一个页表结构

```
struct PageTableEntry {  
    int present; // 指示页面是否在内存中  
    int frame; // 页面的内存帧号  
    // 其他页表项的信息(例如访问权限等)..  
};
```

// 缺页中断处理函数


```
void page_fault_handler() { // 获取引起缺页中断的虚拟内存地址  
    int virtual_address = get_faulting_address(); // 从虚拟地址中解析出页表项中的索引  
    int page_index = get_page_index(virtual_address); // 从页表中获取对应页表项  
    PageTableEntry page_table_entry = get_page_table_entry(page_index);  
    // 判断对应页面是否在内存中  
    if (page_table_entry.present) {  
        // 页面已经在内存中, 可能是因为其他原因引起的中断, 进行相应处理... }  
    else { // 页面不在内存中, 需要将其加载到内存中  
        int frame = allocate_frame(); // 从内存中分配一个合适的页面框  
        page_table_entry.frame = frame; // 更新页表项中的页面框号  
        page_table_entry.present = 1; // 更新页表项中的页面在内存标志  
        load_page_from_disk(page_index, frame); // 从磁盘加载页面到内存中  
    } // 更新页表项中的其他信息...  
    update_page_table_entry(page_index, page_table_entry); // 恢复被中断的指令执行  
    restore_interrupted_instruction();  
}
```



Page Fault (Cont.)

◆ Restart instruction

➤ block move



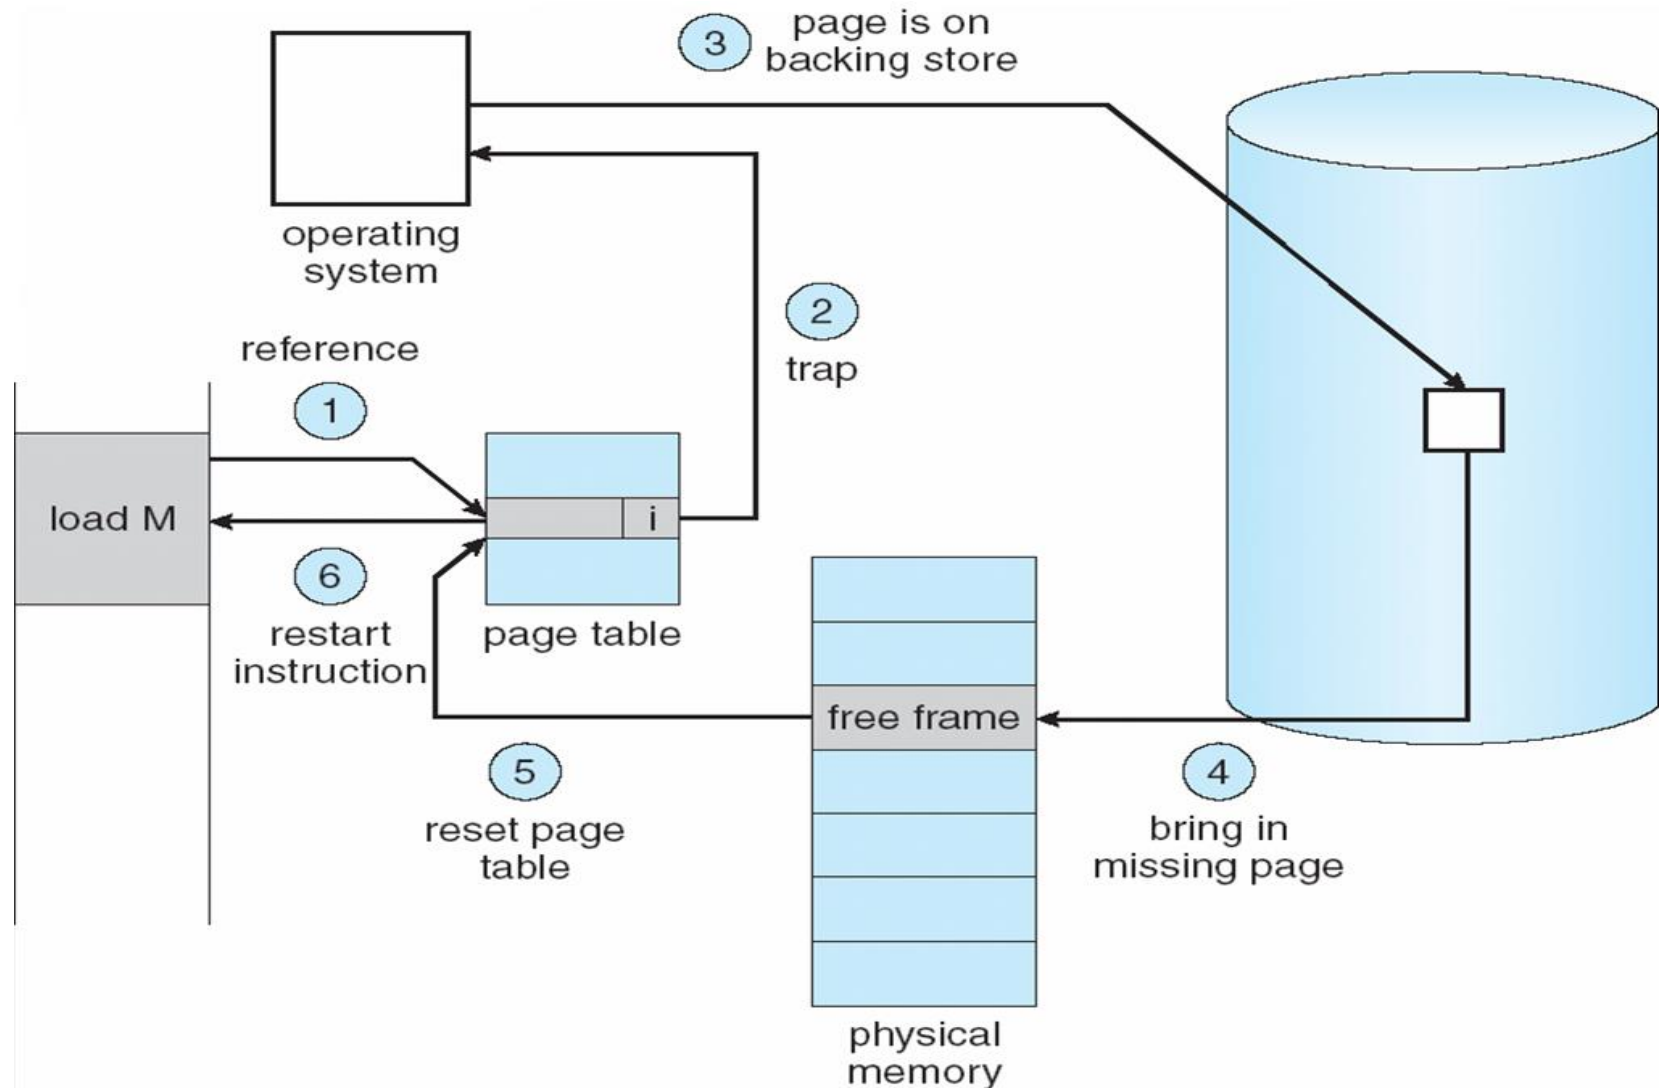
page fault process

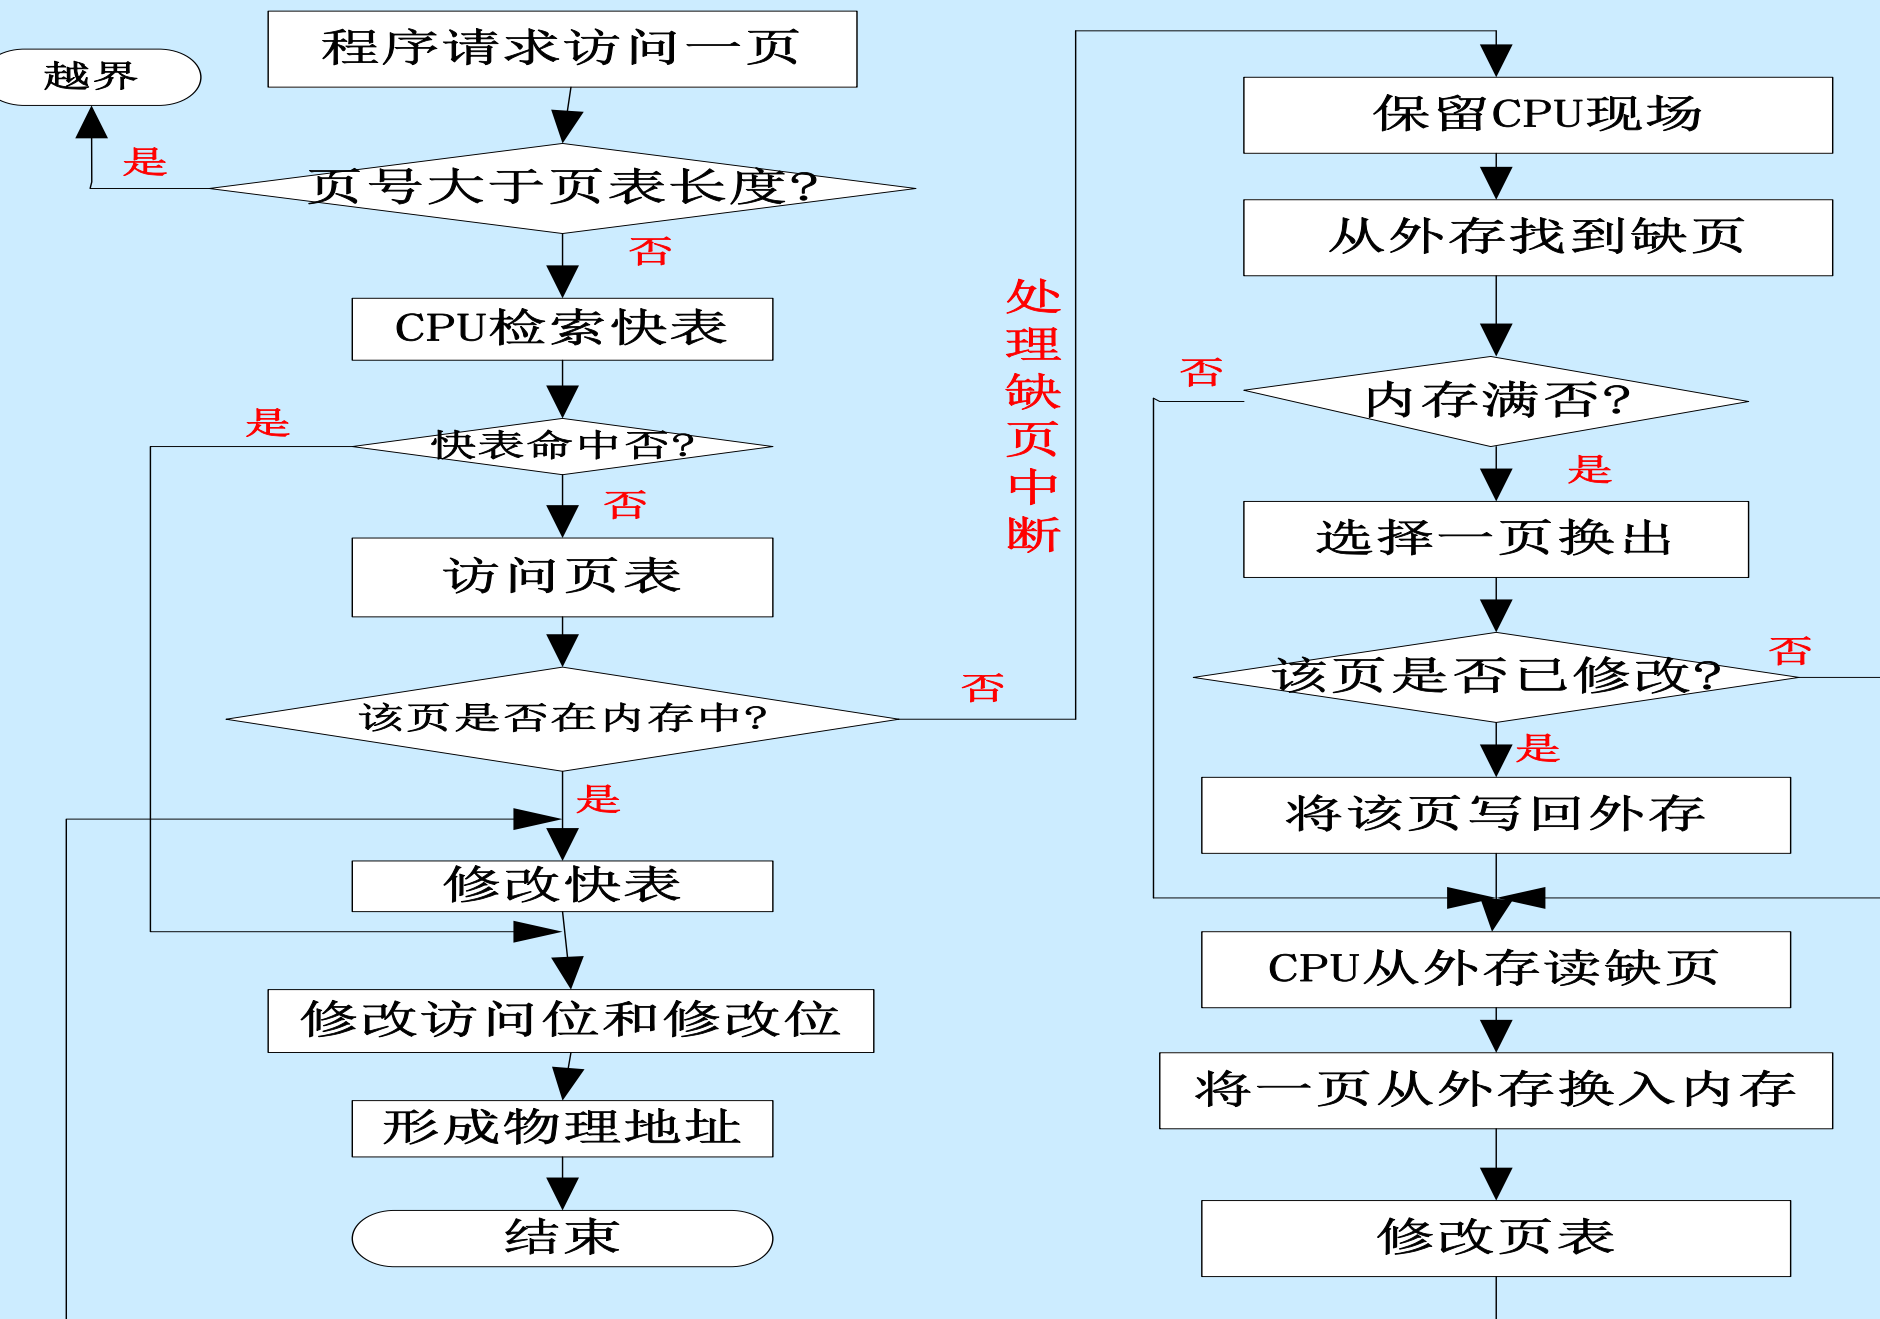
➤ auto increment/decrement location





Steps in Handling a Page Fault





请求分页系统中地址变换的过程



Performance of Demand Paging

- ◆ Page Fault Rate $0 \leq p \leq 1.0$
 - if $p = 0$, no page faults
 - if $p = 1$, every reference is a fault

- ◆ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \text{ (page fault overhead} \\ & + \text{swap page out} \\ & + \text{swap page in} \\ & + \text{restart overhead) } \end{aligned}$$





Demand Paging Example

- ◆ Memory access time = 200 nanoseconds
- ◆ Average page-fault service time = 8 milliseconds(毫秒)
- ◆
$$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$
- ◆ If one access out of 1,000 causes a page fault, then
$$\text{EAT} = 8.2 \text{ microseconds(微妙)}.$$





Process Creation

- ◆ Virtual memory allows other benefits during process creation:
 - Copy-on-Write (写入时复制)
 - Memory-Mapped Files (later)







Copy-on-Write

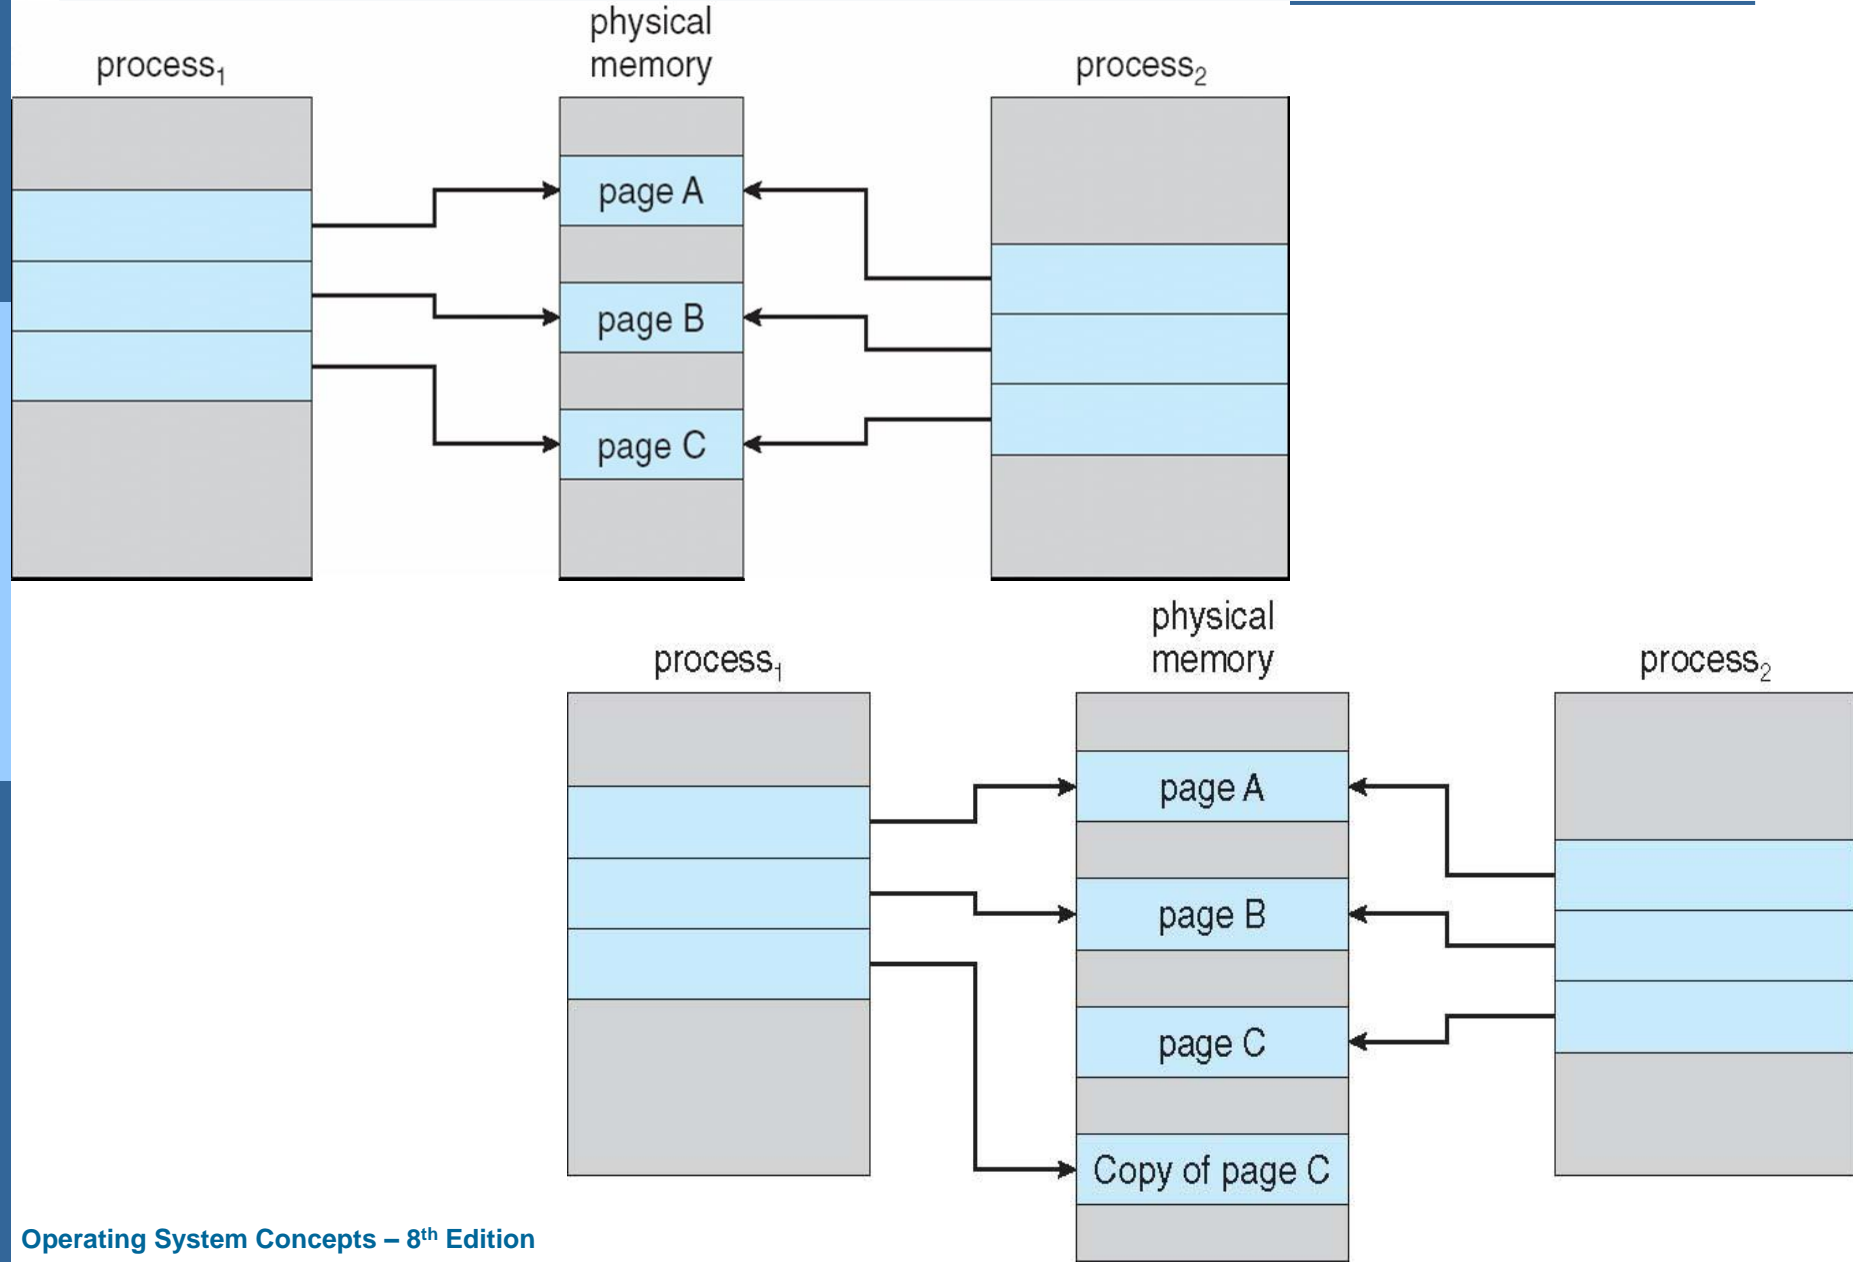
- ◆ Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory. If either process modifies a shared page, only then is the page copied.

在Linux程序中，fork（）会产生一个和父进程完全相同的子进程，出于效率考虑，引入了“写时复制”技术，也就是只有进程空间的各段的内容要发生变化时，才会将父进程的内容复制一份给子进程。





Before Process 1 Modifies Page C





What happens if there is no free frame?

- ◆ Page replacement – find some page in memory, but not really in use, swap it out
 - algorithm
 - performance – want an algorithm which will result in minimum number of page faults
- ◆ Same page may be brought into memory several times







Page Replacement

- ◆ Prevent over-allocation of memory by modifying page-fault service routine to **include page replacement**
- ◆ Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- ◆ Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory





Need For Page Replacement








Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process





Page Replacement





Page Replacement Algorithms


- ◆ Want lowest page-fault rate
- ◆ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
- ◆ In all our examples, the reference string is

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5





Graph of Page Faults Versus The Number of Frames





First-In-First-Out (FIFO) Algorithm

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 $n=12$

◆ 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

◆ 4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

◆ Belady's Anomaly: more frames \Rightarrow more page faults





FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2					2	2	4	4	4	0					0	0			7	7	7
	0	0	0					3	3	3	2	2	2					1	1			1	0	0
		1	1					1	0	0	0	3	3					3	2			2	2	1


page frames

$$F=15 \text{ page faults} \longrightarrow f = F/n = 15/20 = 0.75$$





FIFO Illustrating Belady's Anomaly






Optimal Algorithm

- ◆ Replace page that will not be used for longest period of time
- ◆ 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 $n=12$



- ◆ How do you know this?
- ◆ Used for measuring how well your algorithm performs






Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

$F=9$ page faults

$$f = F/n = 9/20 = 0.45$$





Least Recently Used (LRU) Algorithm

◆ Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

Counter implementation

- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- When a page needs to be changed, look at the counters to determine which are to change





LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

$F=12$ page faults

$f = F/n = 12/20 = 0.6$





LRU Algorithm (Cont.)

- ◆ Stack implementation – keep a stack of page numbers in a double link form: Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
- ✓ No search for replacement

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

2
1
0
7
4

stack
before
a

7
2
1
0
4

stack
after
b

↑ ↑
a b

Use Of A Stack to Record The Most Recent Page References





LRU Approximation Algorithms

◆ Reference bit

- With each page associate a bit, initially = 0
- When page is referenced bit set to 1
- Replace the one which is 0 (if one exists) 替换标示为0的页
 - ▶ We do not know the order, however

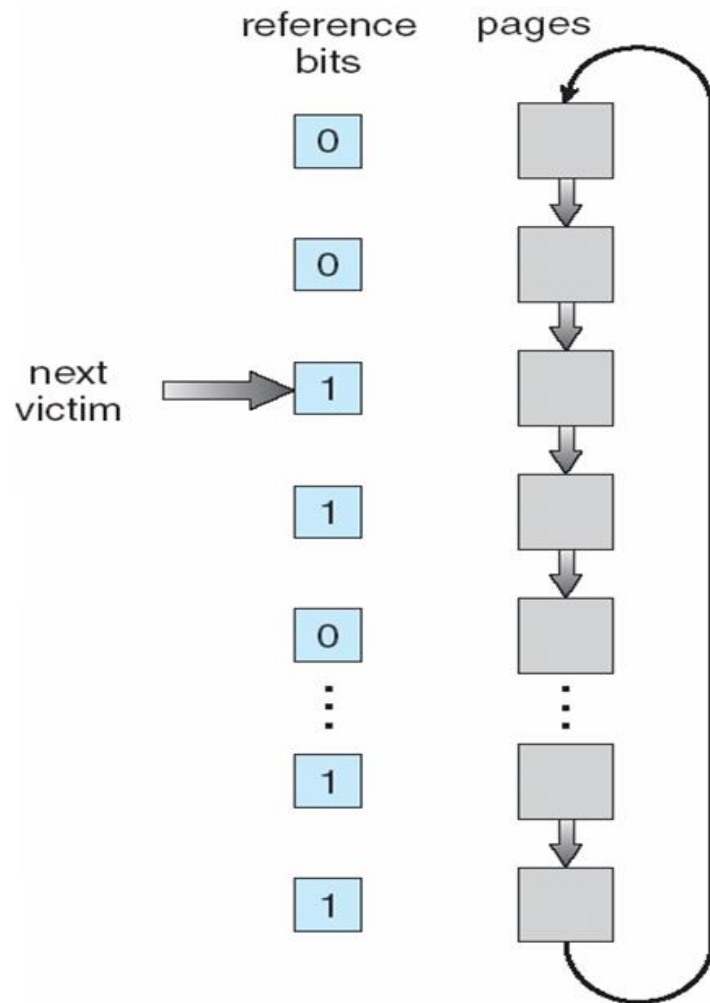
◆ Second chance

- Need reference bit
- Clock replacement
- If page to be replaced (in clock order) has reference bit = 1 then:
(如果被替换的页为1,先置为0, 看看被引用的频率)
 - ▶ set reference bit 0
 - ▶ leave page in memory
 - ▶ replace next page (in clock order), subject to same rules






Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)





Enhanced Second chance Algorithm

- ❑ 不仅考虑页面的使用情况，并考虑置换代价，选择淘汰页面时：选择未访问且未被修改的页面。设置两位
 - ◆ 访问位 (A) ， 修改位 (M) ；
 - ◆ 启动一个进程时，A与M均置为0；
 - ◆ A被周期清零；
- ❑ 内存所有页面分成为四类，选择开销最小的置换。
 - ◆ 第0类：无访问，无修改($A=0, M=0$)
 - ◆ 第1类：无访问，有修改($A=0, M=1$)
 - ◆ 第2类：有访问，无修改($A=1, M=0$)
 - ◆ 第3类：有访问，有修改($A=1, M=1$)





Counting Algorithms

- ◆ Keep a counter of the number of references that have been made to each page
- ◆ **LFU(least frequently used) Algorithm**: replaces page with smallest count (使用最少的)
- ◆ **MFU(Most frequently used) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used (使用最多的)





Allocation of Frames

- ❑ Each process needs *minimum* number of pages
- ❑ Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ◆ instruction is 6 bytes, might span 2 pages
 - ◆ 2 pages to handle *from*
 - ◆ 2 pages to handle *to*
- ❑ Two major allocation schemes
 - ◆ fixed allocation
 - ◆ priority allocation





Fixed Allocation

- ◆ Equal allocation – For example, if there are 100 frames and 5 processes, give each process 20 frames.
- ◆ Proportional allocation – Allocate according to the size of process
 - s_i = size of process p_i
 - $S = \sum s_i$
 - m = total number of frames
 - a_i = allocation for $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$





Priority Allocation

- Use a proportional allocation scheme using priorities rather than size

- If process P_i generates a page fault,
 - ◆ select for replacement one of its frames
 - ◆ select for replacement a frame from a process with lower priority number





Global vs. Local Allocation


- ❑ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- ❑ **Local replacement** – each process selects from only its own set of allocated frames





Thrashing


- ❑ If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - low CPU utilization
 - operating system thinks that it needs to increase the degree of multiprogramming
 - another process added to the system
- ❑ **Thrashing** \equiv a process is busy swapping pages in and out





Demand Paging and Thrashing

- ❑ Why does demand paging work? Locality model
 - ◆ Process migrates from one locality to another
 - ◆ Localities may overlap
- ❑ Why does thrashing occur? Σ size of locality > total memory size



Locality In A Memory-Reference Pattern





Working-Set Model

- ◆ $\Delta \equiv$ working-set window \equiv a fixed number of page references

Example: 10,000 instruction

- ◆ WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ (varies in time)
 - ✓ if Δ too small will not encompass entire locality
 - ✓ if Δ too large will encompass several localities
 - ✓ if $\Delta = \infty \Rightarrow$ will encompass entire program
- ◆ $D = \sum WSS_i \equiv$ total demand frames
- ◆ if $D > m \Rightarrow$ Thrashing
- ◆ Policy if $D > m$, then suspend one of the processes








Working-set model

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...





Keeping Track of the Working Set



- ◆ Approximate with interval timer + a reference bit
- ◆ Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = 1 \Rightarrow page in working set
- ◆ Why is this not completely accurate?
- ◆ Improvement = 10 bits and interrupt every 1000 time units







Page-Fault Frequency Scheme

- ◆ Establish “acceptable” page-fault rate
 - ✓ If actual rate too low, process loses frame
 - ✓ If actual rate too high, process gains frame





Working Sets and Page Fault Rates






Memory-Mapped Files

- ◆ Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory.
- ◆ A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page.
Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- ◆ Simplifies file access by treating file I/O through memory rather than `read()` `write()` system calls.
- ◆ Also allows several processes to map the same file allowing the pages in memory to be shared.






Memory Mapped Files





Memory-Mapped Shared Memory in Windows





Allocating Kernel Memory

- ◆ Treated differently from user memory
- ◆ Often allocated from a free-memory pool
 - Kernel requests memory for structures of varying sizes
 - Some kernel memory needs to be contiguous
- ◆ 分配核心级内存不同于用户级的：大小变化/连续





Design features

- ◆ Which free chunks should service request (分配那个块)
 - Ideally avoid fragmentation... requires future knowledge
- ◆ Split free chunks to satisfy smaller requests
 - Avoids internal fragmentation
- ◆ Coalesce (合并) free blocks to form larger chunks
 - Avoids external fragmentation



- 分区管理问题?





Buddy System

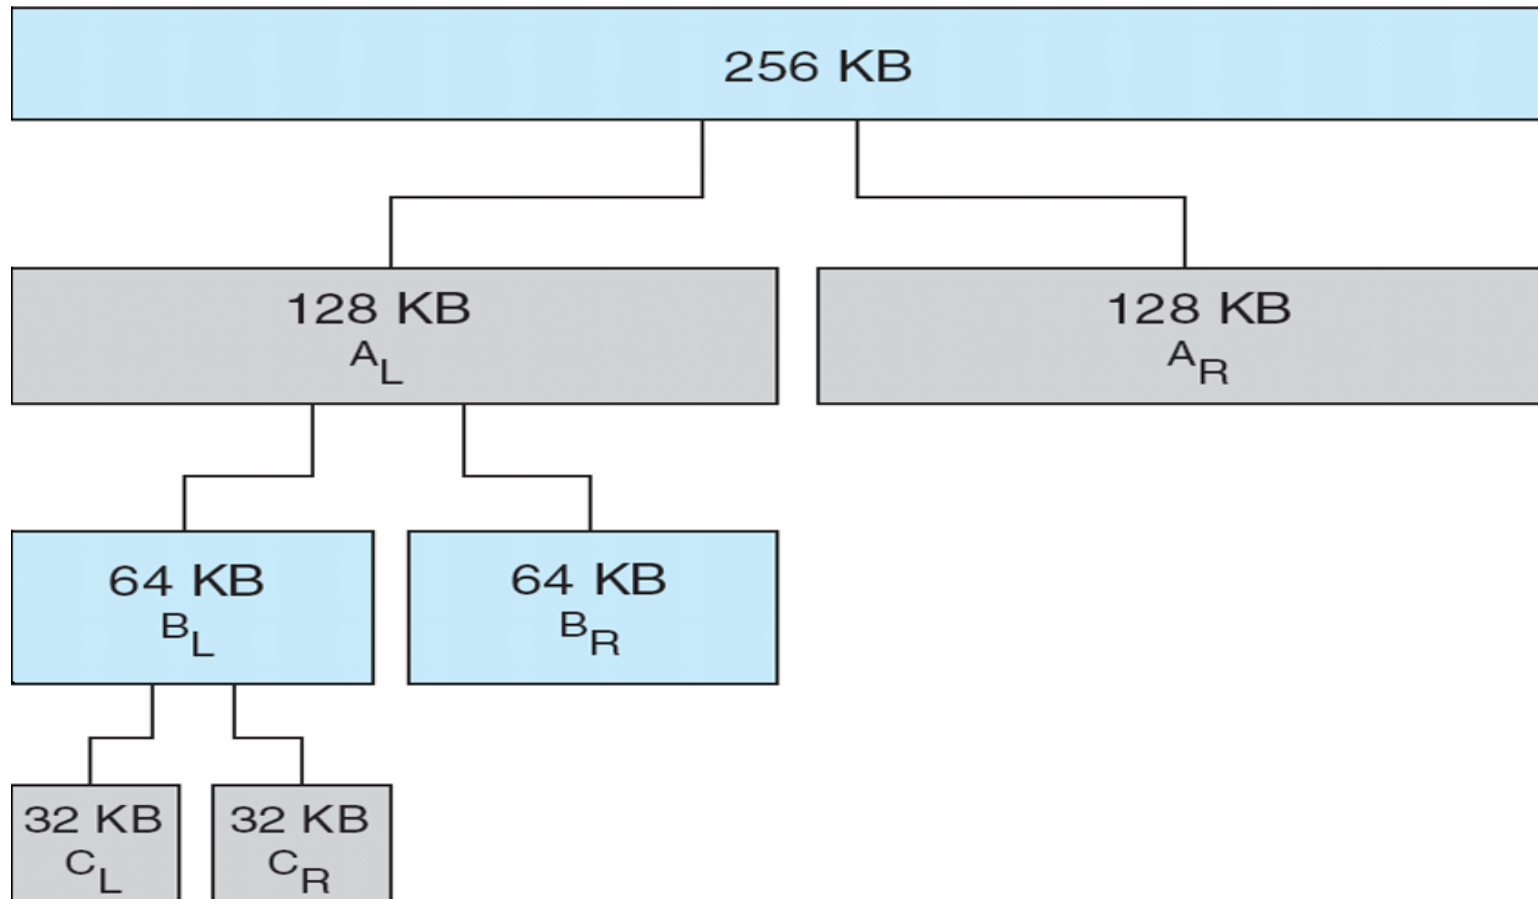
- ◆ Allocates memory from fixed-size segment consisting of physically-contiguous pages
- ◆ Memory allocated using **power-of-2 allocator**
 - ① Satisfies requests in units sized as power of 2
 - ② Request rounded up to next highest power of 2
 - ③ When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - ▶ Continue until appropriate sized chunk available
- ◆ 按2的幂次方大小拆分或合并，分配最接近大小的块。





Buddy System Allocator

physically contiguous pages





Slab Allocator


- ◆ Alternate strategy
- ◆ **Slab** is one or more physically contiguous pages
- ◆ **Cache** consists of one or more slabs (用于cache管理分配)
- ◆ Single cache for each unique kernel data structure
 - ✓ Each cache filled with **objects** – instantiations of the data structure
- ◆ When cache created, filled with objects marked as **free**
- ◆ When structures stored, objects marked as **used**
- ◆ If slab is full of used objects, next object allocated from empty slab
 - ✓ If no empty slabs, new slab allocated
- ◆ Benefits include no fragmentation, fast memory request satisfaction

将分配的内存分割成各种尺寸的块 (chunk)，并把尺寸相同的块分成组 (chunk 的集合)





Slab Allocation





Other Issues -- Prepaging

◆ Prepaging

- ✓ To reduce the large number of page faults that occurs at process startup
- ✓ Prepage all or some of the pages a process will need, before they are referenced
- ✓ But if prepaged pages are unused, I/O and memory was wasted
- ✓ Assume s pages are prepaged and α of the pages is used
 - ▶ Is cost of $s * \alpha$ save pages faults $>$ or $<$ than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages? α near zero \Rightarrow prepaging loses
 - ▶ 缺页中断开销/预分配的页却没有使用的开销 = ?






Other Issues – Page Size

- ◆ Page size selection must take into consideration:
 - ✓ Fragmentation: 页面大，则内碎片大;
 - ✓ table size: 页面小，则页表占用的空间大;
 - ✓ I/O overhead: 磁盘I/O时间中传输时间和数据量有关系;
 - ✓ Locality;






Other Issues – Program Structure

◆ Program structure

- ✓ `Int[128,128] data;`
- ✓ Each row is stored in one page
- ✓ Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i,j] = 0;
```


128 x 128 = 16,384 page faults



✓ Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

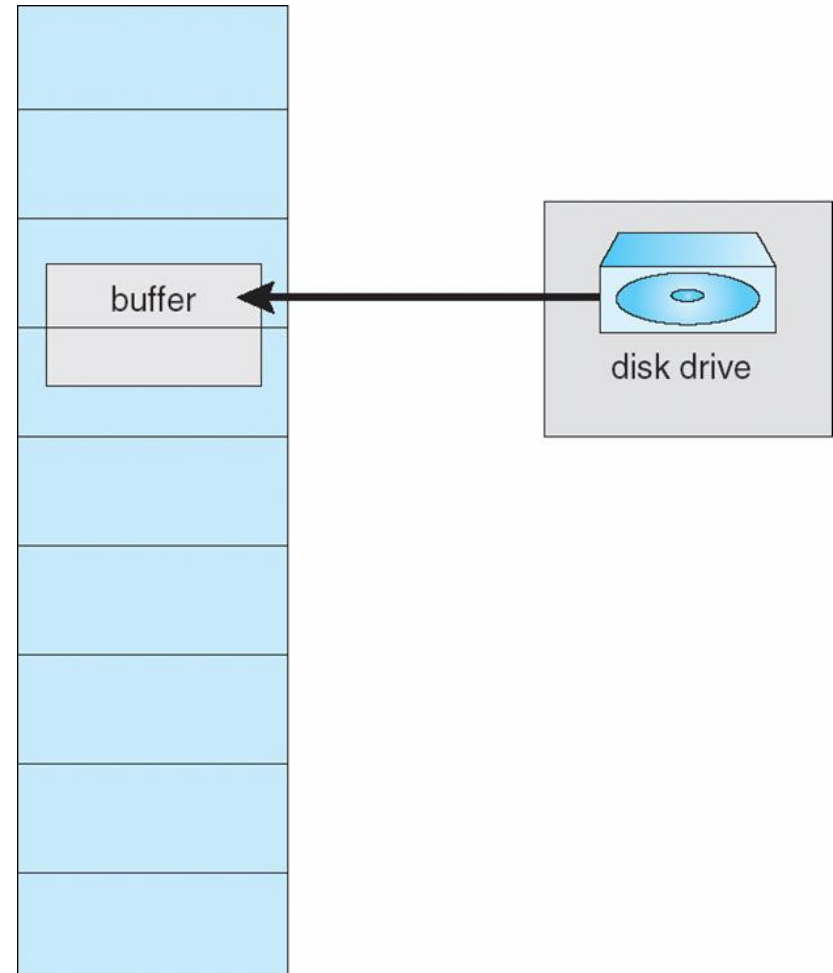
128 page faults





Other Issues – I/O interlock

- ◆ **I/O Interlock** – Pages must sometimes be locked into memory
- ◆ Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm
- ◆ 不能置换




**Reason Why Frames Used For I/O
Must Be In Memory**





Caching




- ◆ **Cache**: a repository for copies that can be accessed more quickly than the original
- ◆ It underlies many of the techniques that are used today to make computers fast (提高性能的技术)
 - ✓ Can cache: memory locations, address translations, pages, file blocks, file names, network routes, etc...
- ◆ Important measure: Average Access time =
 $(\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$



Why Bother with Caching?

Memory Wall Problem

Processor-DRAM Memory Gap

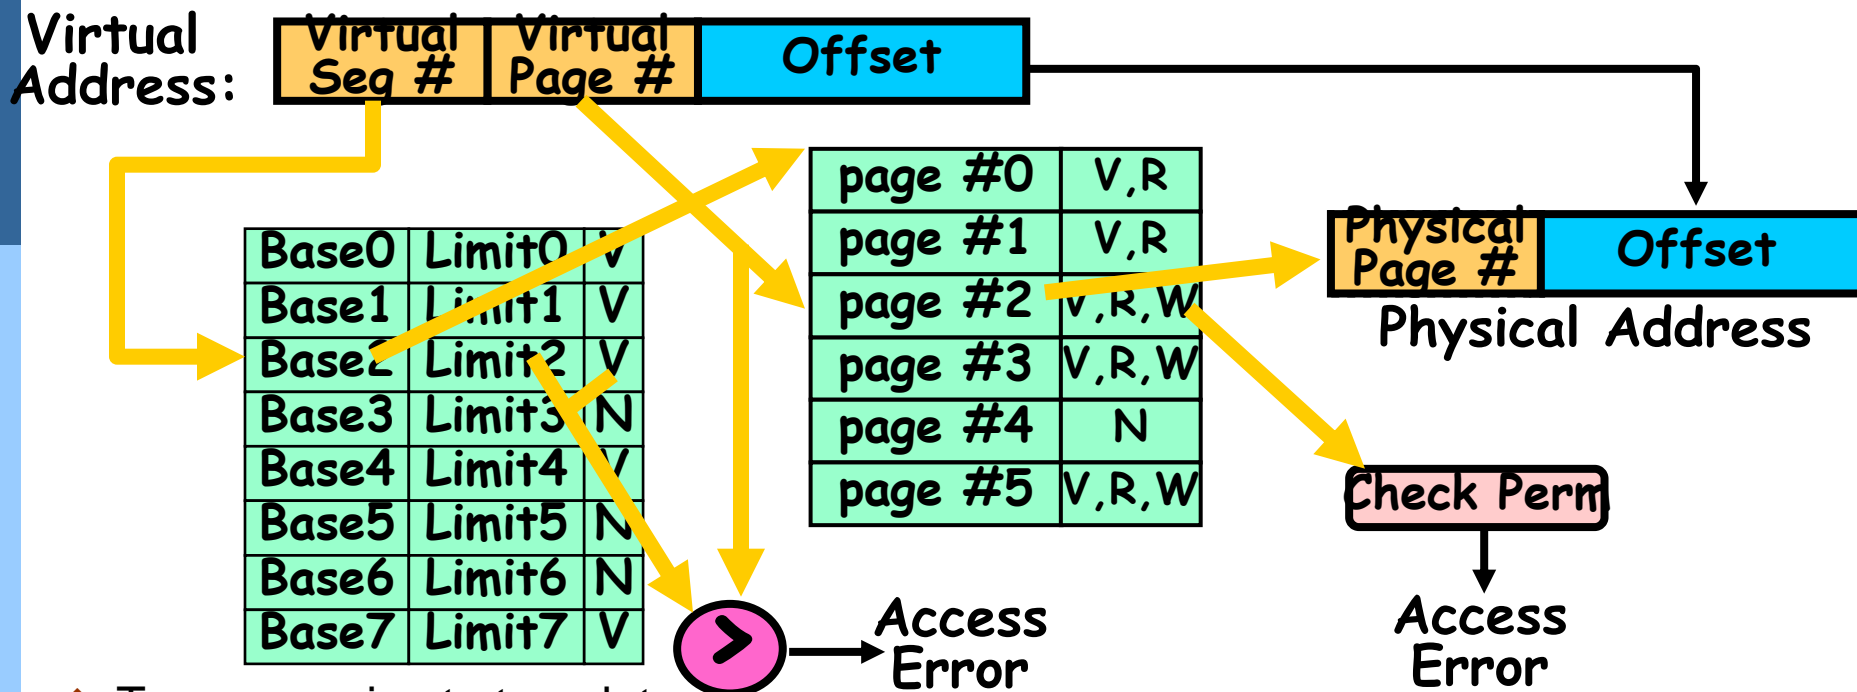


- 1980: no cache in micro-processor; 2010: 3-level cache on chip, 4-level cache off chip
- 1989 the first Intel processor with on-chip L1 cache was Intel 486, 8KB size
- 1995 the first Intel processor with on-chip L2 cache was Intel Pentium Pro, 256KB size
- 2003 the first Intel processor with on-chip L3 cache was Intel Itanium 2, 6MB size

Xian-He Sun



Another Major Reason to Deal with Caching

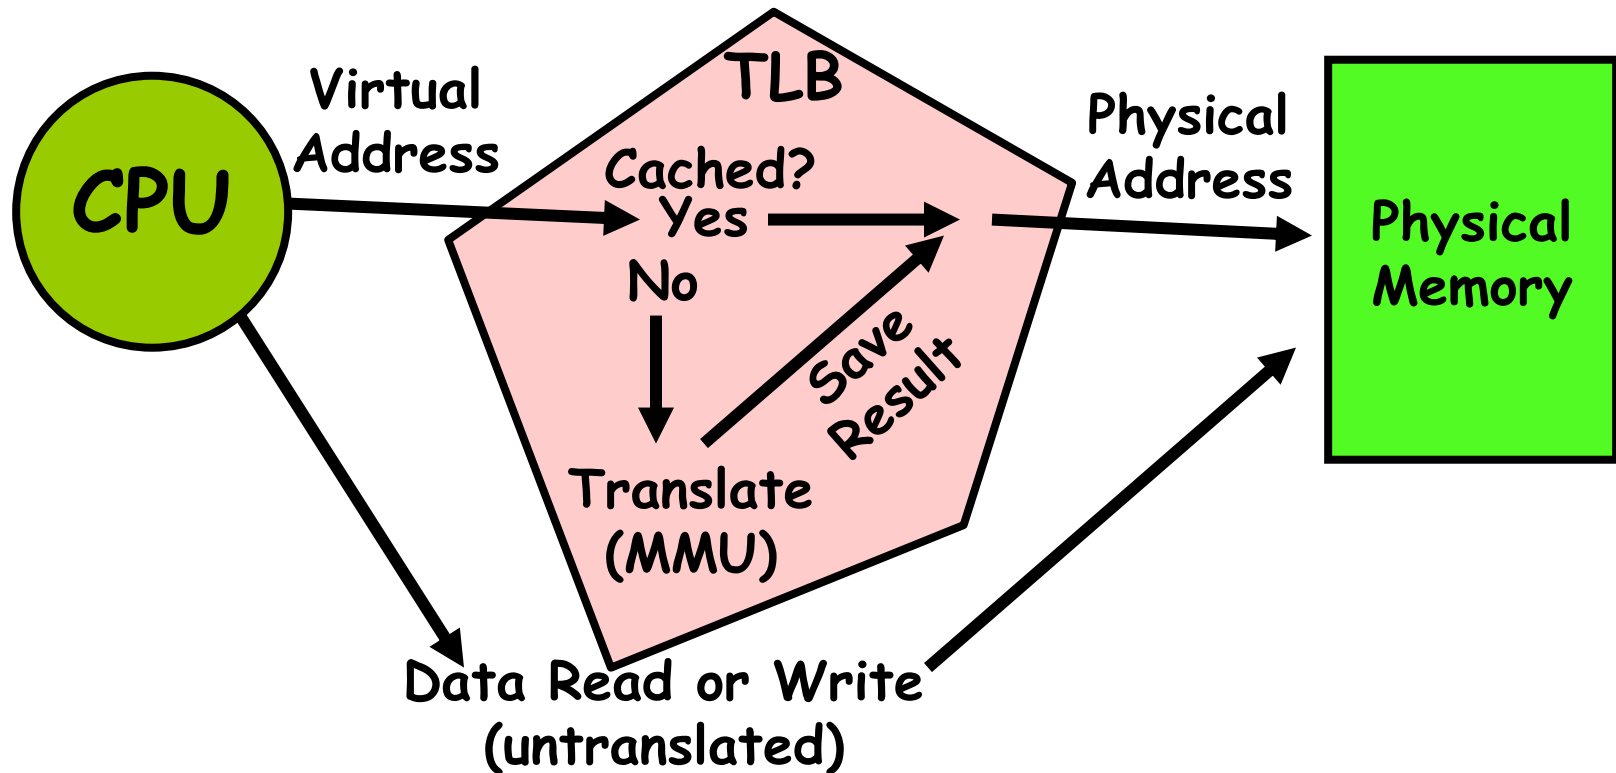


- ◆ Too expensive to translate on every access
 - ✓ At least two DRAM accesses per actual DRAM access
 - ✓ Or: perhaps I/O if page table partially on disk!
- ◆ Even worse problem: What if we are using caching to make memory access faster than DRAM access???
- ◆ Solution? Cache translations!
 - ✓ **Translation Cache: TLB ("Translation Lookaside Buffer")**





Caching Applied to Address Translation





TLB Reach

- ◆ TLB Reach - The amount of memory accessible from the TLB
- ◆ $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- ◆ Ideally, the working set of each process is stored in the TLB
 - ✓ Otherwise there is a high degree of page faults
- ◆ Increase the Page Size
 - ✓ This may lead to an increase in fragmentation as not all applications require a large page size (增加页尺寸产生内碎片)
- ◆ Provide Multiple Page Sizes
 - ✓ This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation







Why Does Caching Help? Locality!



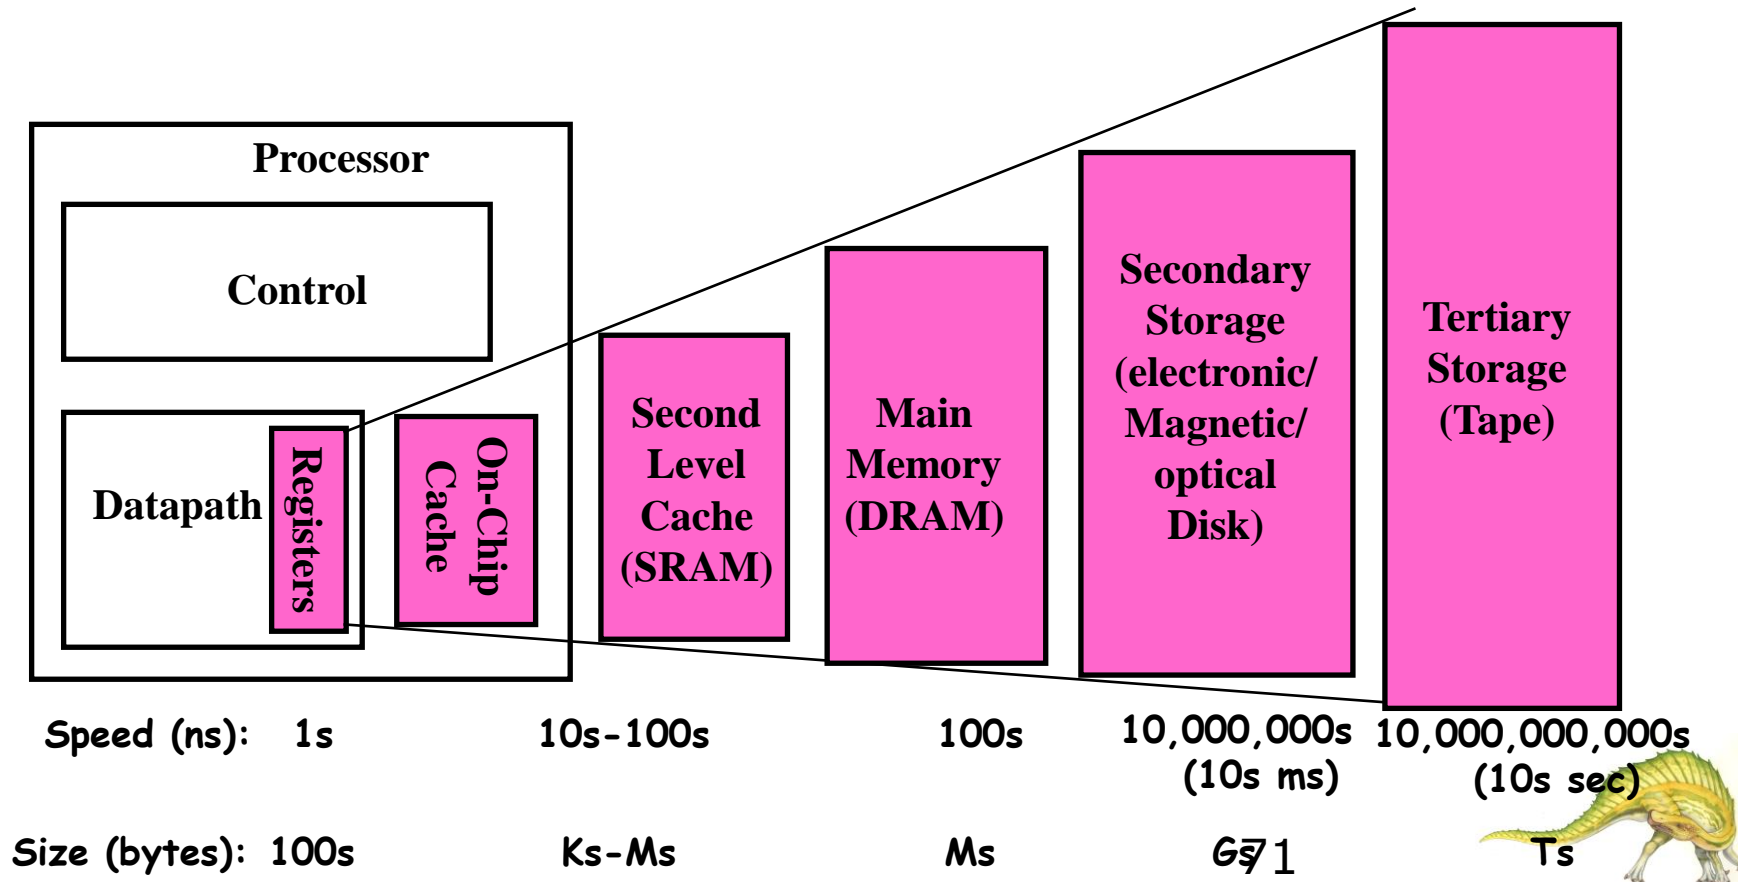
- ◆ **Temporal Locality** (Locality in Time): 最近访问的数据项更接近处理器
- ◆ **Spatial Locality** (Locality in Space): 相邻的数据项更接近处理器





Review: Memory Hierarchy of a Modern Computer System

- ◆ Take advantage of the principle of locality to:
 - ✓ Present as much memory as in the cheapest technology
 - ✓ Provide access at speed offered by the fastest technology





Operating System Examples

◆ Windows XP

◆ Solaris





Windows XP

- ◆ Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page
- ◆ Processes are assigned **working set minimum** and **working set maximum**

Working set minimum \leq

A process may be assigned as many pages

\leq working set maximum

- ◆ When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- ◆ Working set trimming removes pages from processes that have pages in excess of their working set minimum





Solaris

- ◆ Maintains a list of free pages to assign faulting processes
 - **Lotsfree** – threshold parameter (amount of free memory) to begin paging (空闲页数阈值, 大于它则进行分页, 小于它则执行回收)
 - **Desfree** – threshold parameter to increasing paging (最小空闲页阈值, 小于它则加快页面回收)
 - **Minfree** – threshold parameter to being swapping Paging is performed by **pageout** process, **Pageout** scans pages using modified clock algorithm (最小空闲页极限阈值, 小于它则每次请求都进行页面回收)
 - **Scanrate** is the rate at which pages are scanned. This ranges from *slowscan* to *fastscan* (扫描的频率)
- ◆ **Pageout** is called more frequently depending upon the amount of free memory available





Solaris

- ◆ *Lotsfree* – threshold parameter (amount of free memory) to begin paging
 - *Lotsfree* is typically set to 1/64 the size of the physical memory
 - 4 times per second, the kernel checks whether the amount of free memory is less than *lotsfree*
 - ▶ If below *lotsfree*, a process **pageout** starts up, which is similar to the second-chance algorithm
- ◆ If free memory falls below **desfree**, **pageout** will run 100 times per second with the intention of keeping at least **desfree** free memory available






Solaris

- ◆ If the **pageout** is unable to keep free memory at **desfree** for a **30-second average**, the kernel begins swapping processes.
 - Free all pages allocated to swapped processes.
 - The kernel looks for processes that have been idle for a long time.
- ◆ If the system unable to maintain free memory at **minfree**, the **pageout** process is called for every request for a new page.





Solaris 2 Page Scanner





Solaris Virtual Memory Layers

Global Page Replacement Manager — Page Scanner

Address Space Management

segkmem
Kernel Memory
Segment

segmap
File Cache Memory
Segment

segvn
Process Memory
Segment

Hardware Address Translation (HAT) Layer

sun4c
HAT layer

sun4m
HAT layer

sun4d
HAT layer

sun4u
HAT layer

x86
HAT layer

sun4c
sun4-mmu

32/32-bit
4K pages

sun4m
sr-mmu

32/36-bit
4K pages

sun4d
sr-mmu

32/36-bit
4K pages

sun4u
sf-mmu

64/64-bit
8K/4M pages

x86
i386 mmu

32/36-bit
4K pages





assignment

◆ 9.4

◆ 9.8

- ◆ 编程题：编写一个程序，实现本章所述的FIFO、LRU和最优页面置换算法。首先，生成一个随机的页面引用串，页码范围为0~9. 将这个随机页面引用串应用到每个算法，并记录每个算法引起的缺页错误的数量。实现置换算法，以便页面帧的数量可以从1~7。假设采用请求分页。

