


Chapter 11: File System Implementation





Chapter 11: File System Implementation

- ◆ File-System Structure
- ◆ File-System Implementation
- ◆ Directory Implementation
- ◆ Allocation Methods
- ◆ Free-Space Management
- ◆ Efficiency and Performance
- ◆ Recovery
- ◆ NFS
- ◆ Example: WAFL File System





Objectives

- ◆ To describe the details of implementing local file systems and directory structures
- ◆ To describe the implementation of remote file systems
- ◆ To discuss block allocation and free-block algorithms and trade-offs





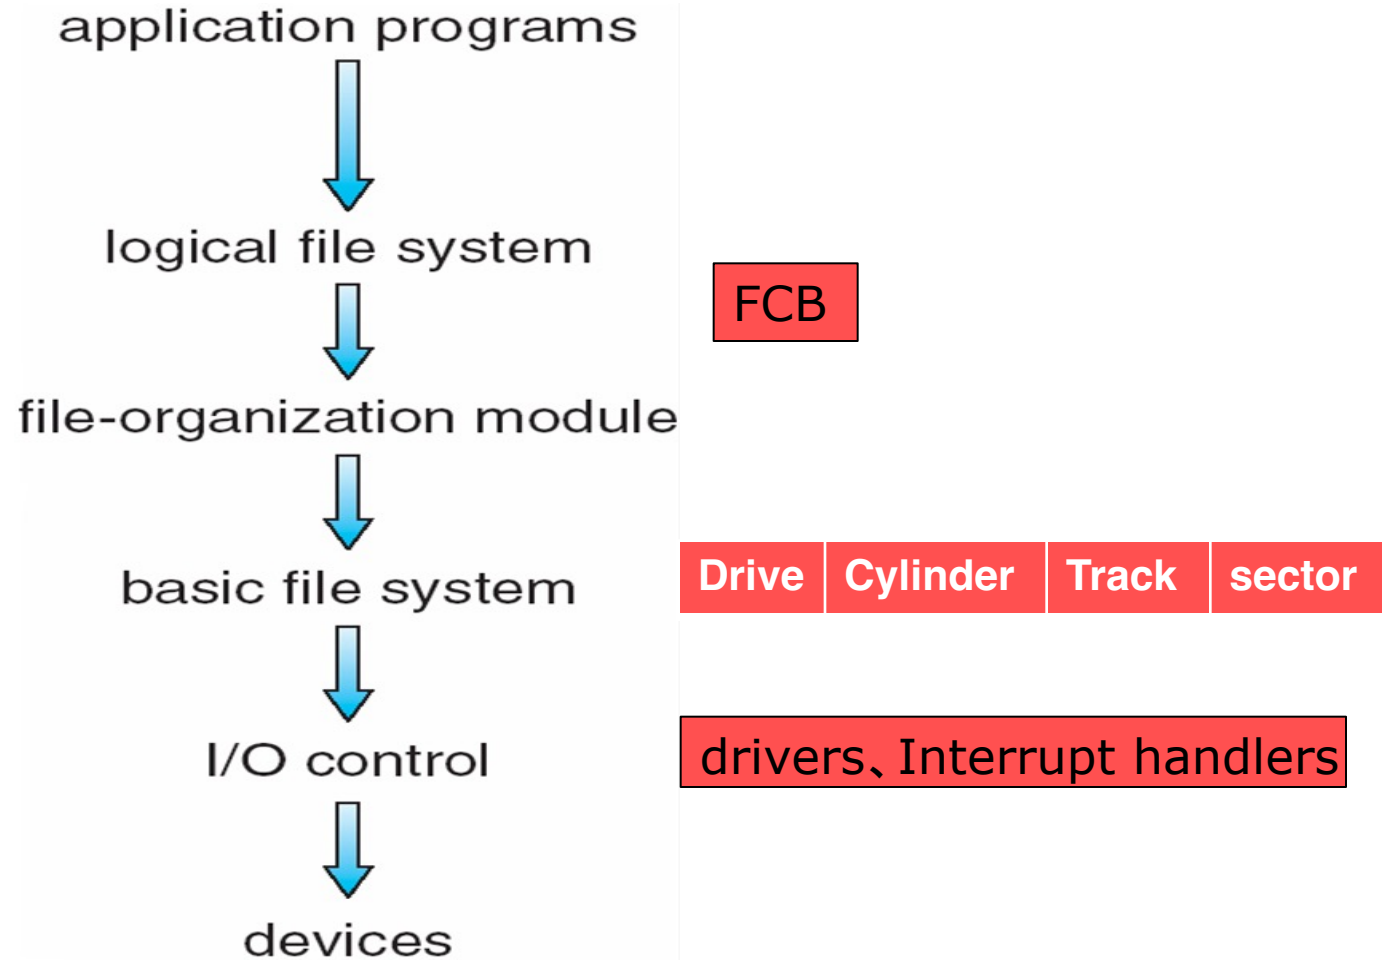
File-System Structure

- ◆ File structure
 - Logical storage unit; Collection of related information
- ◆ File system organized into layers
- ◆ File system resides on secondary storage (disks)
 - Provides efficient and convenient access to disk by allowing data to be stored, located retrieved easily
- ◆ File control block – storage structure consisting of information about a file
- ◆ Device driver controls the physical device





Layered File System





File-System Implementation

- ◆ **Boot control block** contains info needed by system to boot OS from that volume
- ◆ **Volume control block** contains volume details
- ◆ Directory structure organizes the files
- ◆ Per-file **File Control Block (FCB)** contains many details about the file

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

A Typical File Control Block





Implementing File Operations

◆ Create a file:

- Find space in the file system, add directory entry.

◆ Open file

- System call specifying name of file. system searches directory structure to find file.
- 根据文件号查系统打开文件表，看文件是否已被打开；
 - ✓ 是→共享计数加1
 - ✓ 否→将外存中的FCB等信息填入系统打开文件表空表项，共享计数置为1；
- 在用户打开文件表中取一空表项，填写打开方式等，并指向系统打开文件表对应表项
- System keeps **current file position pointer** to the location where next write/read occurs
- System call returns **file descriptor** (a handle) to user process

◆ Writing in a file:

- System call specifying file descriptor and information to be written
- Writes information at location pointed by the files current pointer





Implementing File Operations

- ◆ Reading a file:
 - System call specifying file descriptor and number of bytes to read (and possibly where in memory to stick contents).
- ◆ Repositioning within a file:
 - System call specifying file descriptor and new location of current pointer
 - (also called a file **seek** even though does not interact with disk)
- ◆ Closing a file:
 - System call specifying file descriptor
 - Call removes current file position pointer and file descriptor associated with process and file(打开文件表)
 - 若在文件打开期间，该文件作过某种修改，则应将其写回到辅存。
- ◆ Deleting a file:
 - Search directory structure for named file, release associated file space and erase directory entry
- ◆ Truncating a file:
 - Keep attributes the same, but reset file size to 0, and reclaim file space.





Multiple users of a file

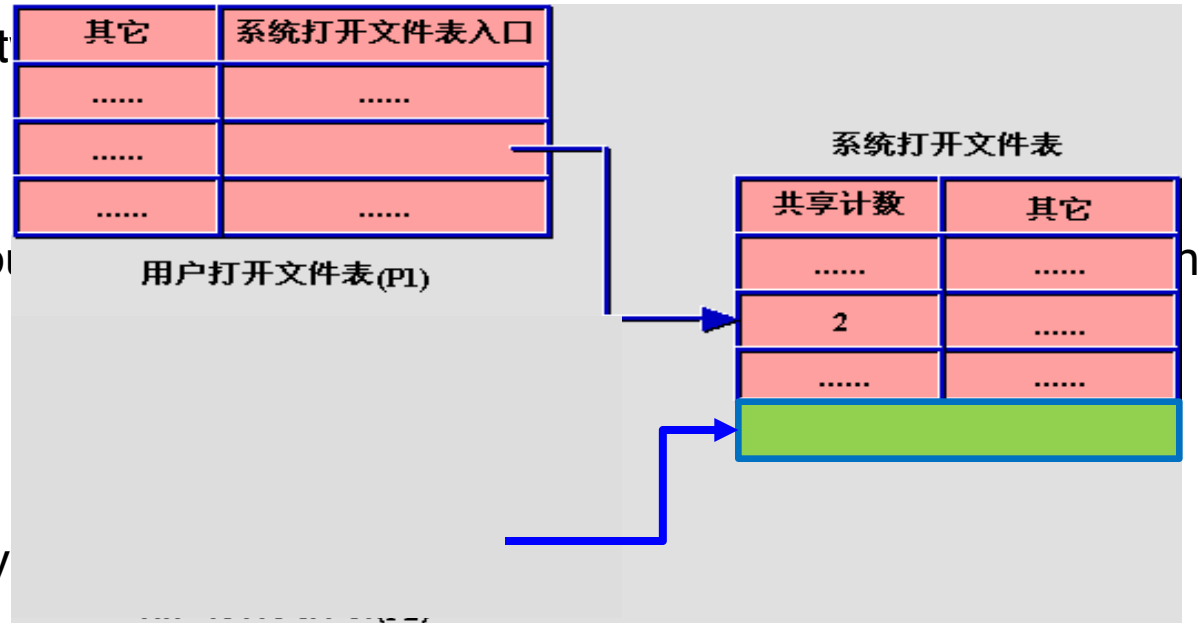
- ◆ OS typically keeps two tables

- ◆ Per-process table

- Information about the file (e.g., file pointer)


- ◆ System wide table

- Gets created by the kernel
- Location of file on disk
- Access dates
- File size
- Count of how many processes have the file open (used for deletion)

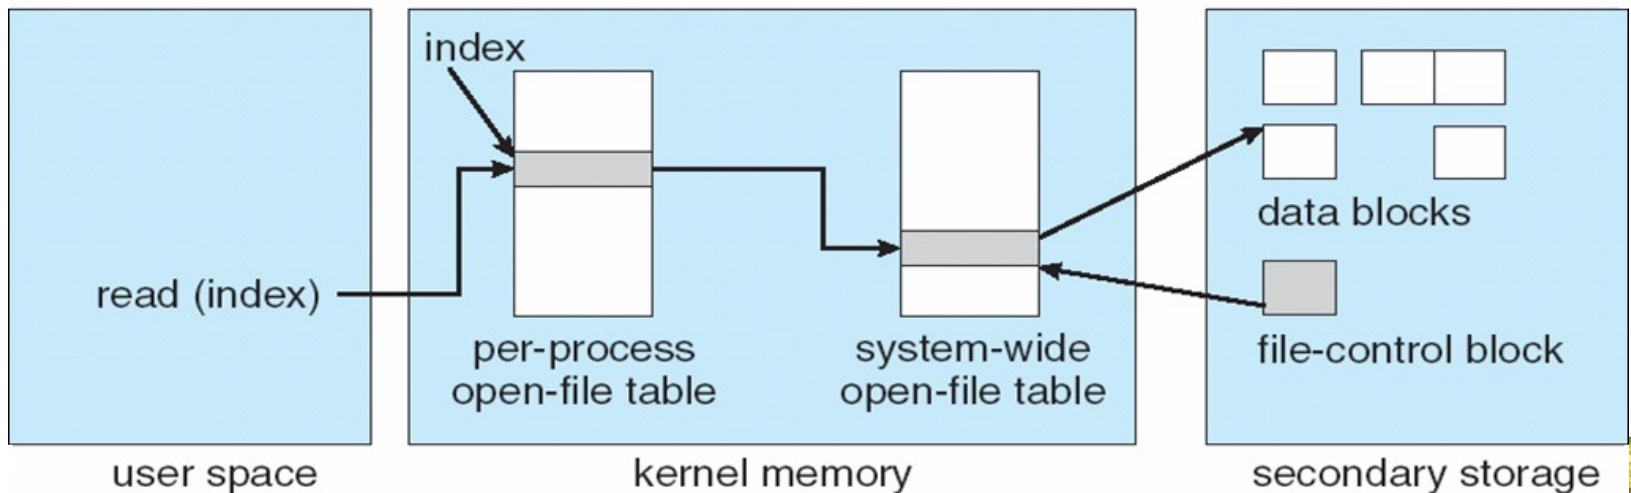




In-Memory File System Structures



(a)




(b)





Virtual File Systems

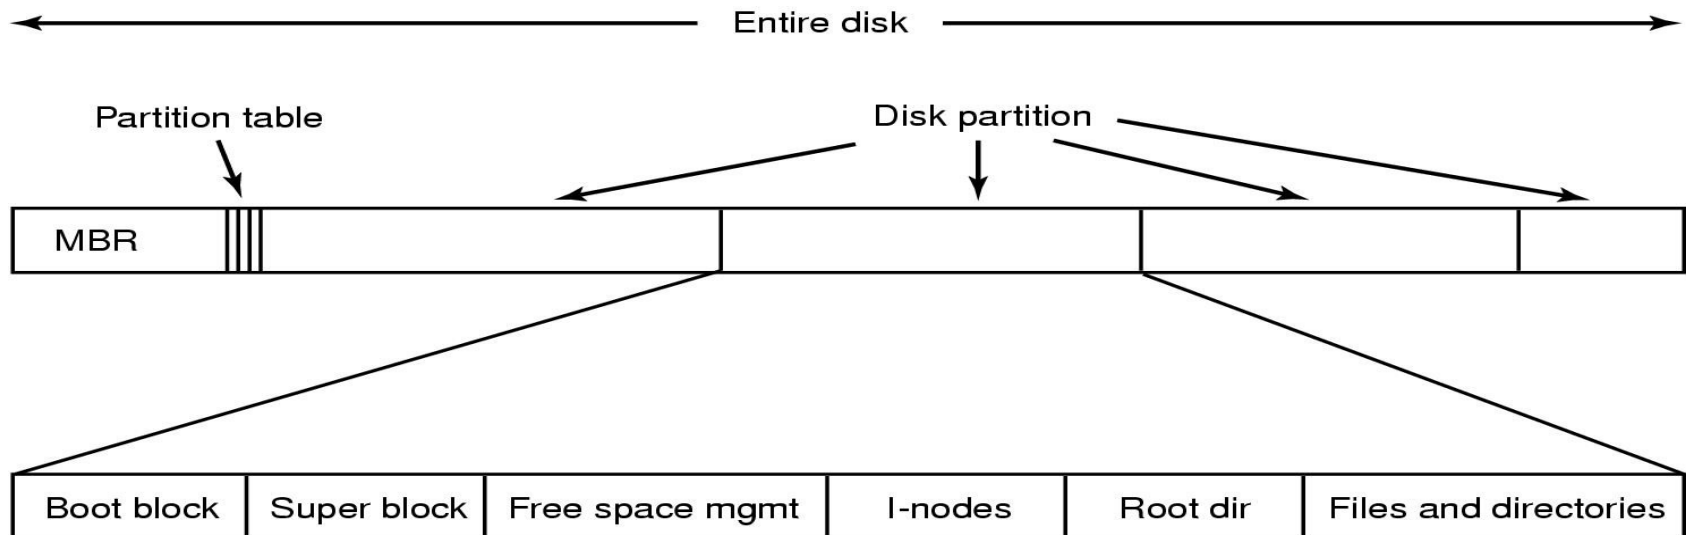
- ◆ Virtual File Systems (VFS) provide an object-oriented way of implementing file systems. VFS allows the same system call interface (the API) to be used for different types of file systems.
- ◆ The API is to the VFS interface, rather than any specific type of file system.





File System Layout

- ◆ File System is stored on disks
 - Disk is divided into 1 or more partitions
 - Sector 0 of disk called Master Boot Record
 - End of MBR has partition table (start & end address of partitions)
- ◆ First block of each partition has boot block
 - Loaded by MBR and executed on boot





Storing Files

Files can be allocated in different ways:

- ◆ Contiguous allocation

- All bytes together, in order

- ◆ Linked Structure

- Each block points to the next block

- ◆ Indexed Structure

- An index block contains pointer to many other blocks

Rhetorical Questions -- which is best?


- For sequential access? Random access?
- Large files? Small files? Mixed?






Implementing Files

- ◆ Contiguous Allocation: allocate files contiguously on disk




(a)



(b)

- ◆ Each file occupies a set of contiguous blocks on the disk;
- ◆ Simple – only starting location (block #) and length (number of blocks) are required; Random access;
- ◆ Wasteful of space (dynamic storage-allocation problem); Files cannot grow,






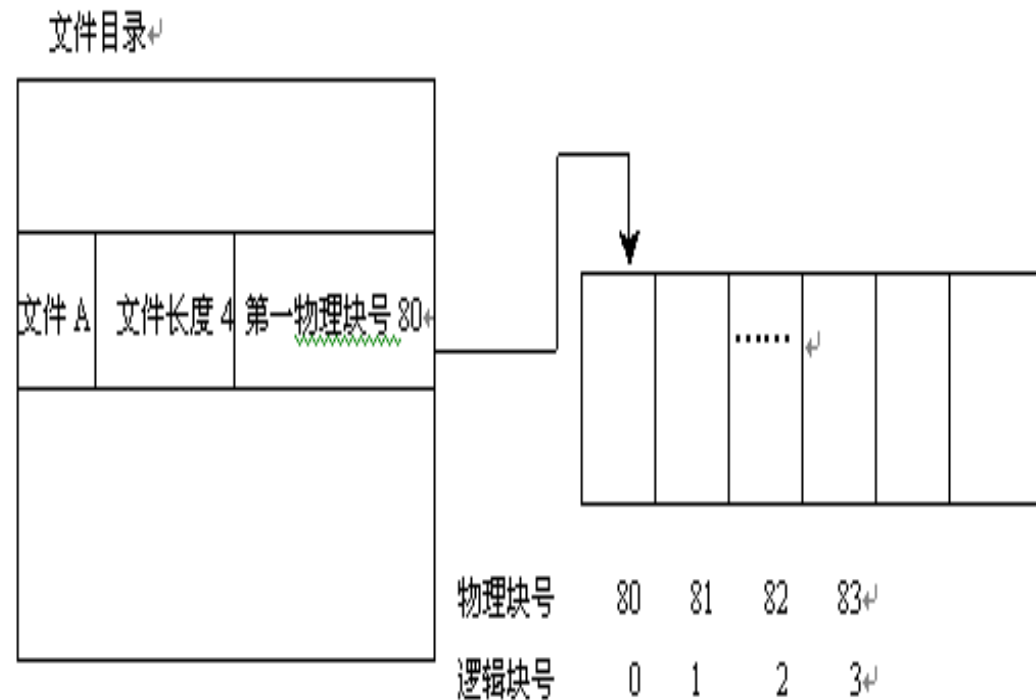
Contiguous Allocation

◆ Mapping from logical to physical

LA/512
Q块
R位移

Block to be accessed = $Q + \text{starting address}$

Displacement into block = R



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

**Contiguous
Allocation of Disk
Space**



Contiguous Allocation

◆ Pros:

- Simple: state required per file is start block and size
- Performance: entire file can be read with one seek

◆ Cons:

- Fragmentation: external is bigger problem
- Usability: user needs to know size of file

◆ Used in CDRoms, DVDs





Extent-Based Systems

- ◆ Many newer file systems (I.e. Veritas File System) use a modified contiguous allocation scheme
- ◆ Extent-based file systems allocate disk blocks in extents
- ◆ An **extent** is a contiguous block of disks
 - Extents are allocated for file allocation
 - A file consists of one or more extents






Linked Allocation

- ◆ Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk. block =

pointer

- ◆ Simple – need only starting address
- ◆ Free-space management system – no waste of space
- ◆ No random access
- ◆ Mapping

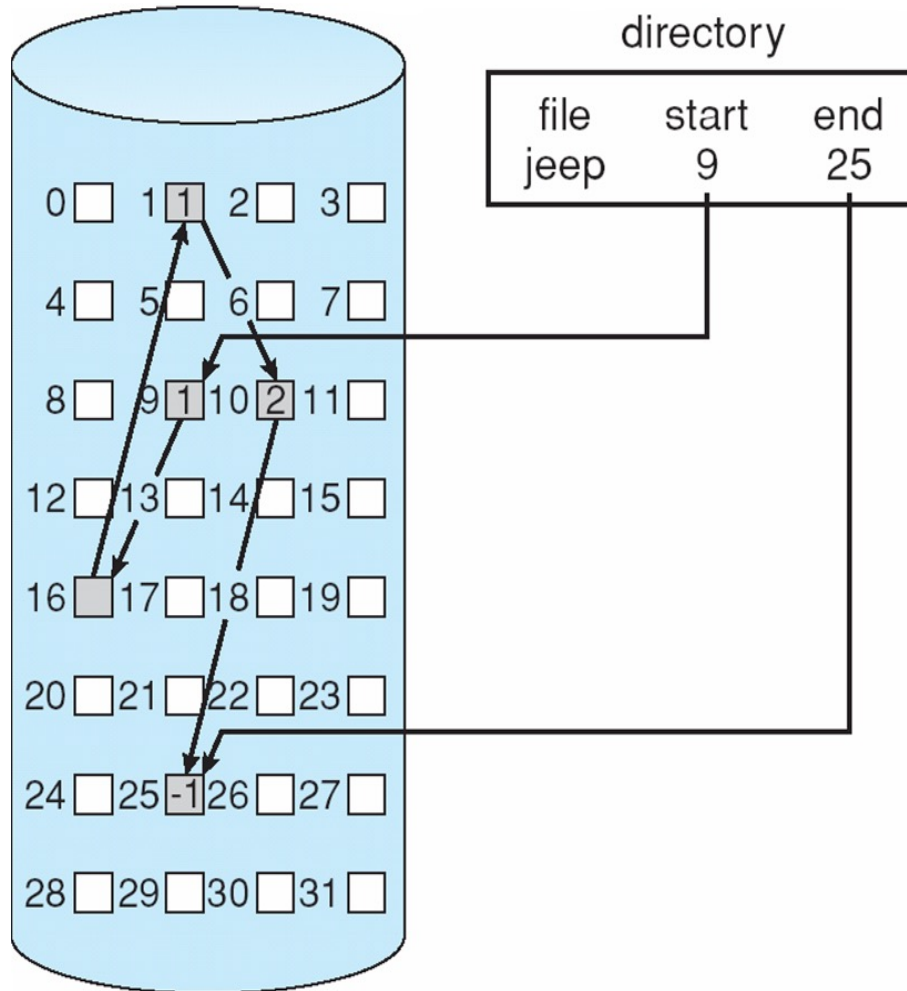


- ◆ Block to be accessed is the Qth block in the linked chain of blocks representing the file.
- ◆ Displacement into block = $R + 1$ (指针占一字节)





Linked Allocation



◆ Pros:

- No space lost to external fragmentation
- Disk only needs to maintain first block of each file

◆ Cons:

- Random access is costly
- Overheads of pointers






MS-DOS(OS/2) File System



- ◆ Implement a linked list allocation using a table
 - Called File Allocation Table (FAT)
 - Take pointer away from blocks, store in this table
 - Can cache FAT in-memory

directory entry

test	...	217
name		start block



File-Allocation Table FAT



Physical block

0	
1	
2	10
3	11
4	7
5	
6	3
7	2
8	
9	
10	12
11	14
12	-1
13	
14	-1
15	

← File A starts here

← File B starts here

← Unused block



FAT Discussion

◆ Pros:

- Entire block is available for data
- Random access is faster since entire FAT is in memory

◆ Cons:


- Entire FAT should be in memory
 - For 20 GB disk, 1 KB block size, FAT has 20 million entries
 - If 4 bytes used per entry \Rightarrow 80 MB of main memory required for FS






Indexed Allocation

- ◆ Brings all pointers together into the **index block**
- ◆ Logical view




Example





Indexed Allocation (Cont.)


- ◆ Need index table
- ◆ Random access
- ◆ Dynamic access without external fragmentation, but have overhead of index block
- ◆ Mapping from logical to physical in a file of maximum size of 256K words and block size of 512 words. We need only 1 block for index table



Q = displacement into index table

R = displacement into block





Indexed Allocation – Mapping (Cont.)

- ◆ Mapping from logical to physical in a file of unbounded length (block size of 512 words)

扩充的实现方案如下：

1. Linked scheme – Link blocks of index table (no limit on size)

$$LA / (512 \times 511) \begin{cases} Q_1 \\ R_1 \end{cases}$$


Q_1 = block of index table


R_1 is used as follows:

$$R_1 / 512 \begin{cases} Q_2 \\ R_2 \end{cases}$$

Q_2 = displacement into block of index table

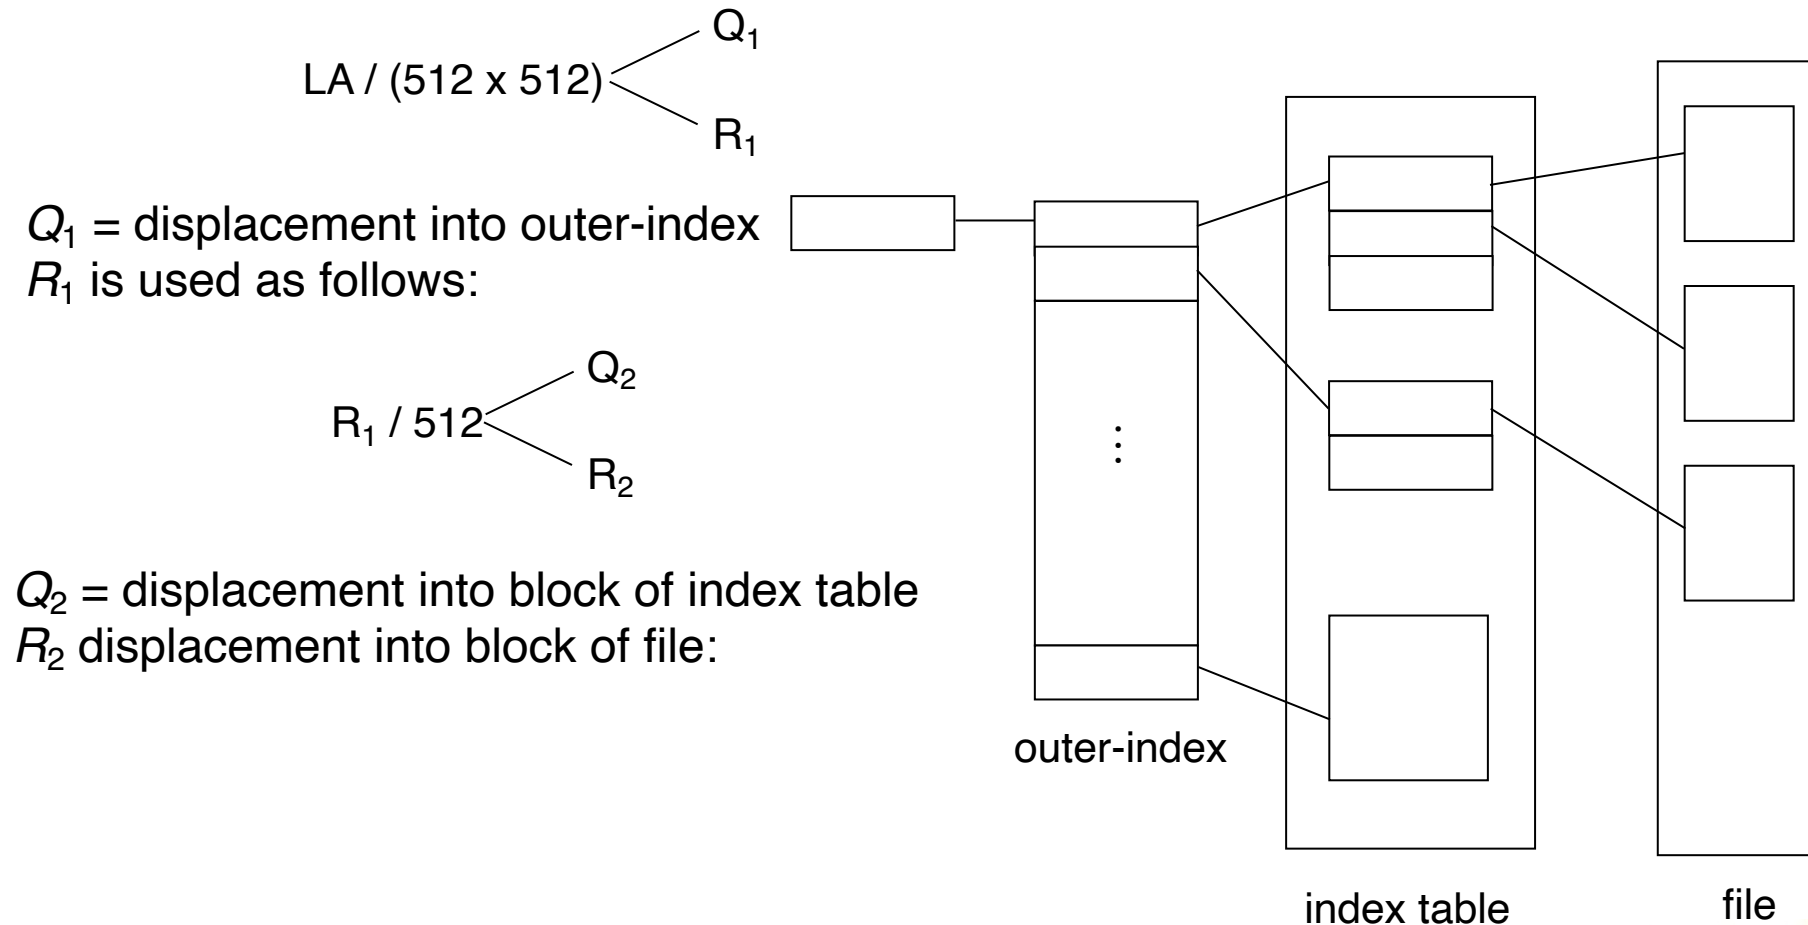
R_2 displacement into block of file:





Indexed Allocation – Mapping (Cont.)

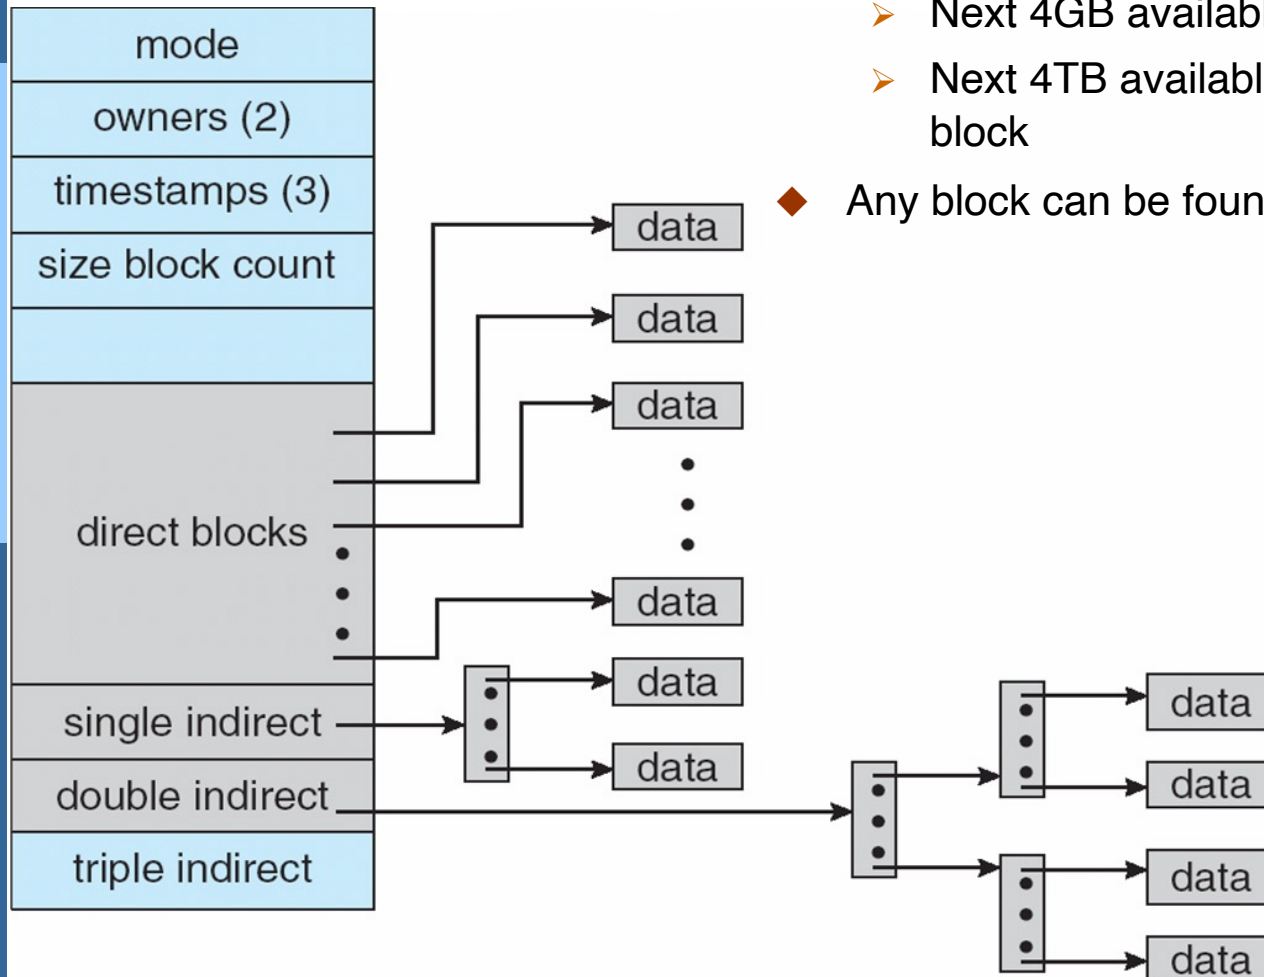
2. Multilevel (for exam. Two-level index)index (maximum file size is 512^3)





3. Combined Scheme: UNIX UFS (4K bytes per block)

共15个指针项，前
12个是直接指针



- ◆ If data blocks are 4K ...
 - First 48K reachable from the inode
 - Next 4MB available from single-indirect
 - Next 4GB available from double-indirect
 - Next 4TB available through the triple-indirect block
- ◆ Any block can be found with at most 3 disk accesses





Questions?

- ◆ Performance?
- ◆ Efficiency?
 - For sequential access? Random access?
 - Large files? Small files? Mixed?
 - HDD VS. SSD?







Managing Free Disk Space

2 approaches to keep track of free disk blocks

◆ Linked list and bitmap approach



(a)




(b)





Free-Space Management



$$\text{bit}[i] = \begin{cases} 0 \Rightarrow \text{block}[i] \text{ free} \\ 1 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

Block number calculation

(number of bits per word) * (number of 0-value words) + offset of first 1 bit

◆ Bit map requires extra space

➤ Example:

◆ block size = 2^{12} bytes

◆ disk size = 2^{30} bytes (1 gigabyte)

◆ $n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)

◆ Easy to get contiguous files





Linked Free Space List on Disk

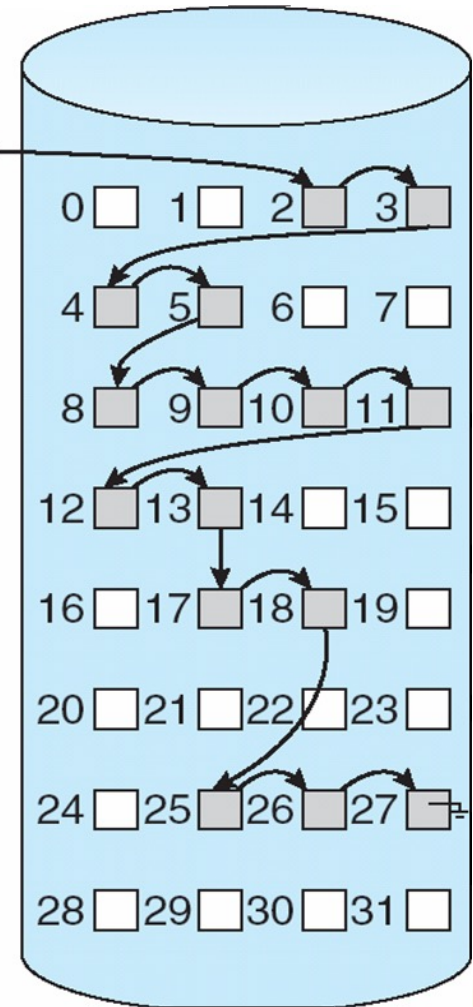
◆ Linked list (free list)

- Cannot get contiguous space easily
- No waste of space

空闲链的改进方法：

- ◆ Grouping(空闲块按组存放)
- ◆ Counting(利用连续性：只记录第一个块地址和后续连续的空闲块个数)

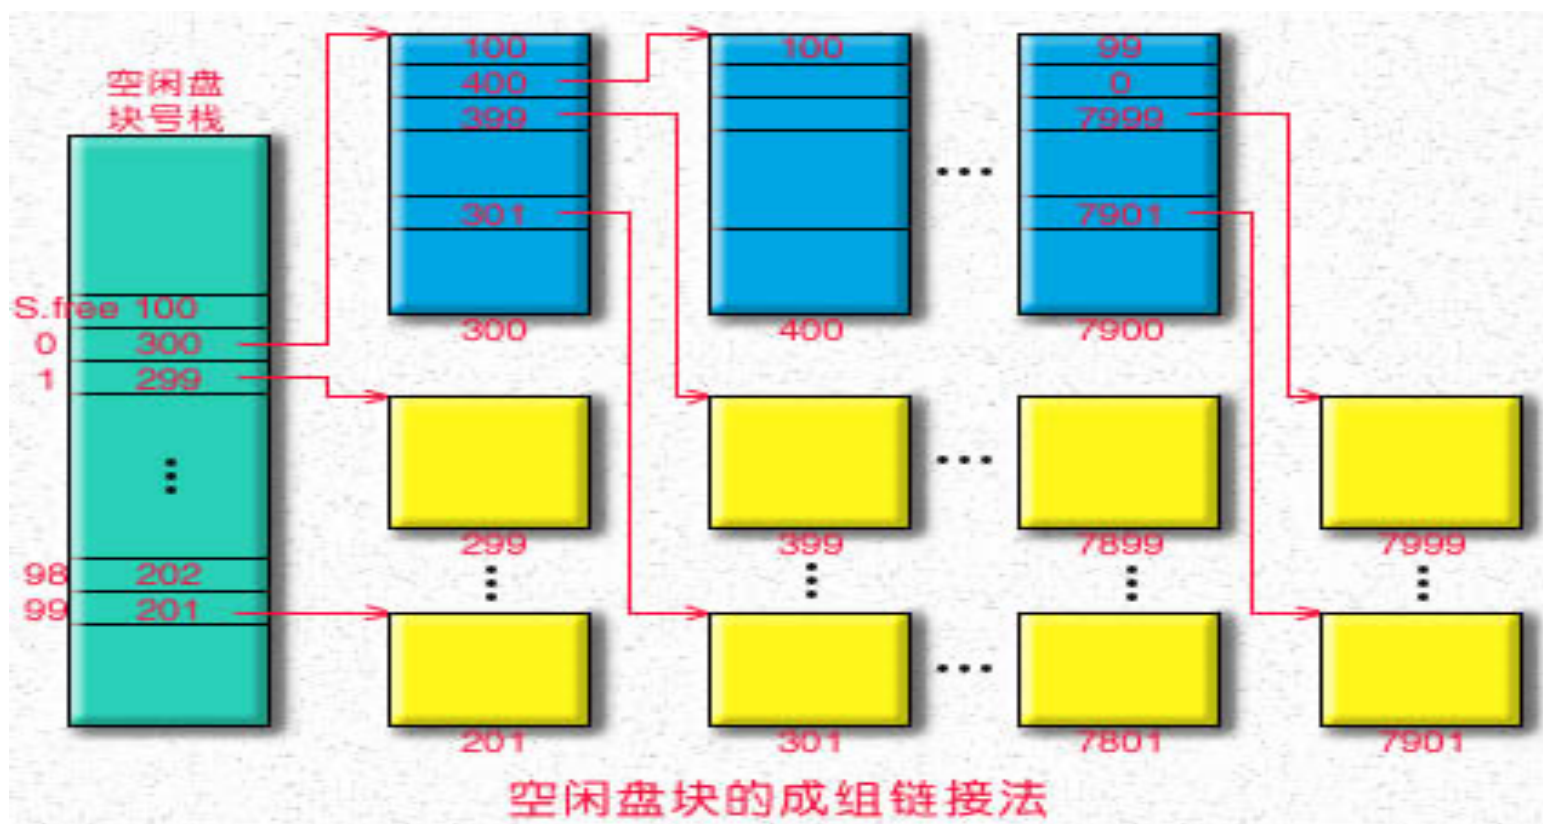
free-space list head





UNIX的成组链接法

- ◆ **空闲盘块号栈**：存放当前可用的空闲盘块号及空闲盘块号数N（最多100个）；文件区的所有空闲盘块被分成若干组；
- ◆ **第一组盘块总数和盘块号**记入空闲盘块号栈中；每一组的第一个盘块中记录有下一组的盘块总数和盘块号






Directory Implementation

- ◆ When a file is opened, OS uses path name to find dir
 - Directory has information about the file's disk blocks
 - Whole file (contiguous), first block (linked-list) or I-node
 - Directory also has attributes of each file
- ◆ Directory: **map ASCII file name to file attributes & location**
- ◆ 2 options: entries have all attributes, or point to file I-node

games	attributes
mail	attributes
news	attributes
work	attributes

(a)



(b)

Data structure
containing the
attributes





UNIX采用(b), 给定文件路径名为 /usr / ast / mbox, 检索过程如下:

根目录

1	.
1	..
4	bin
7	dev
14	lib
9	etc
6	usr
8	tmp

结点6是
/usr的目录

132

132#块是/usr
的目录

6	.
1	..
19	disk
30	erik
51	jim
26	ast
45	bat

结点26是
/usr/ast
目录

406

406#块是
/usr/ast的
目录

26	.
6	..
64	grants
92	books
60	mbox
81	minix
17	src





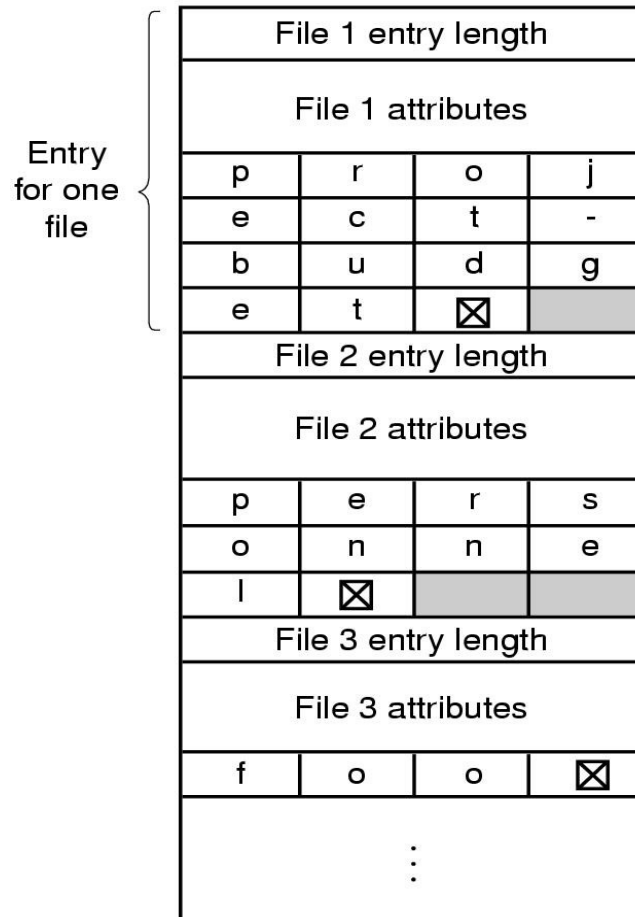
Directory Implementation

- ◆ Linear list of file names with pointer to the data blocks
 - simple to program
 - time-consuming to execute
- ◆ Hash Table – linear list with hash data structure
 - decreases directory search time
 - **collisions** – situations where two file names hash to the same location
 - fixed size

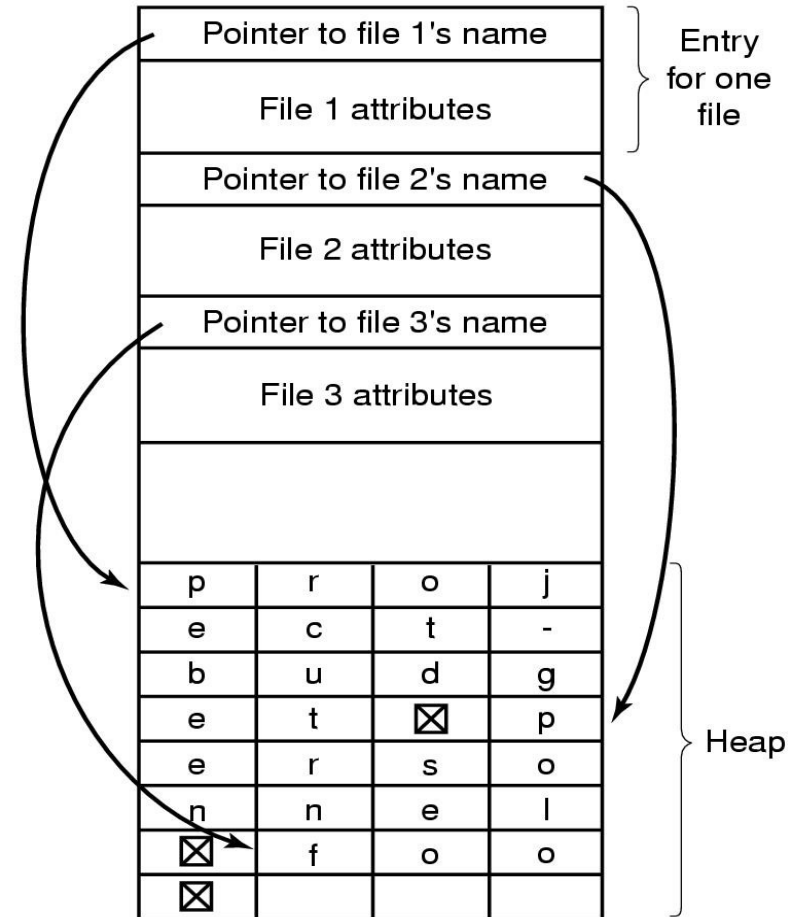




Managing file names: Example



(a)



(b)





Efficiency and Performance

◆ Efficiency dependent on:

- disk allocation and directory algorithms
- types of data kept in file's directory entry

◆ Performance

- disk cache – separate section of main memory for frequently used blocks
- free-behind and read-ahead – techniques to optimize sequential access (预先读/延迟写)
- improve PC performance by dedicating section of memory as virtual disk, or RAM disk

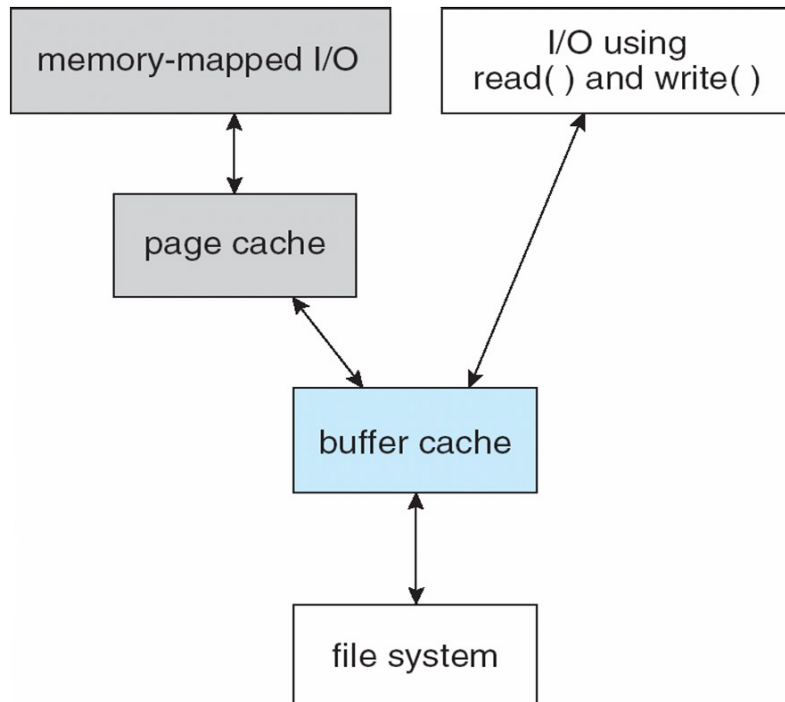





Page Cache

- ◆ A **page cache** caches pages rather than disk blocks using virtual memory techniques
- ◆ Memory-mapped I/O uses a page cache
- ◆ Routine I/O through the file system uses the buffer (disk) cache
- ◆ This leads to the following figure :

I/O Without a Unified Buffer Cache



I/O Using a Unified Buffer Cache





Recovery

- ◆ **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
- ◆ Use system programs to **back up** data from disk to another storage device (magnetic tape, other magnetic disk, optical)
- ◆ Recover lost file or disk by **restoring** data from backup





Log Structured File Systems

- ◆ Log structured (or journaling) file systems record each update to the file system as a transaction
- ◆ All transactions are written to a log
 - A transaction is considered committed once it is written to the log
 - However, the file system may not yet be updated
- ◆ The transactions in the log are asynchronously written to the file system
 - When the file system is modified, the transaction is removed from the log
- ◆ If the file system crashes, all remaining transactions in the log must still be performed
- ◆ Example.

```
BEGIN_TRANSACTION;
```

```
x = x + 1;
```

```
y = y + 2
```

```
x = y * y;
```

```
END_TRANSACTION;
```

Log

[x = 0 / 1]

[y = 0/2]

[x = 1/4]



写完后系统崩溃





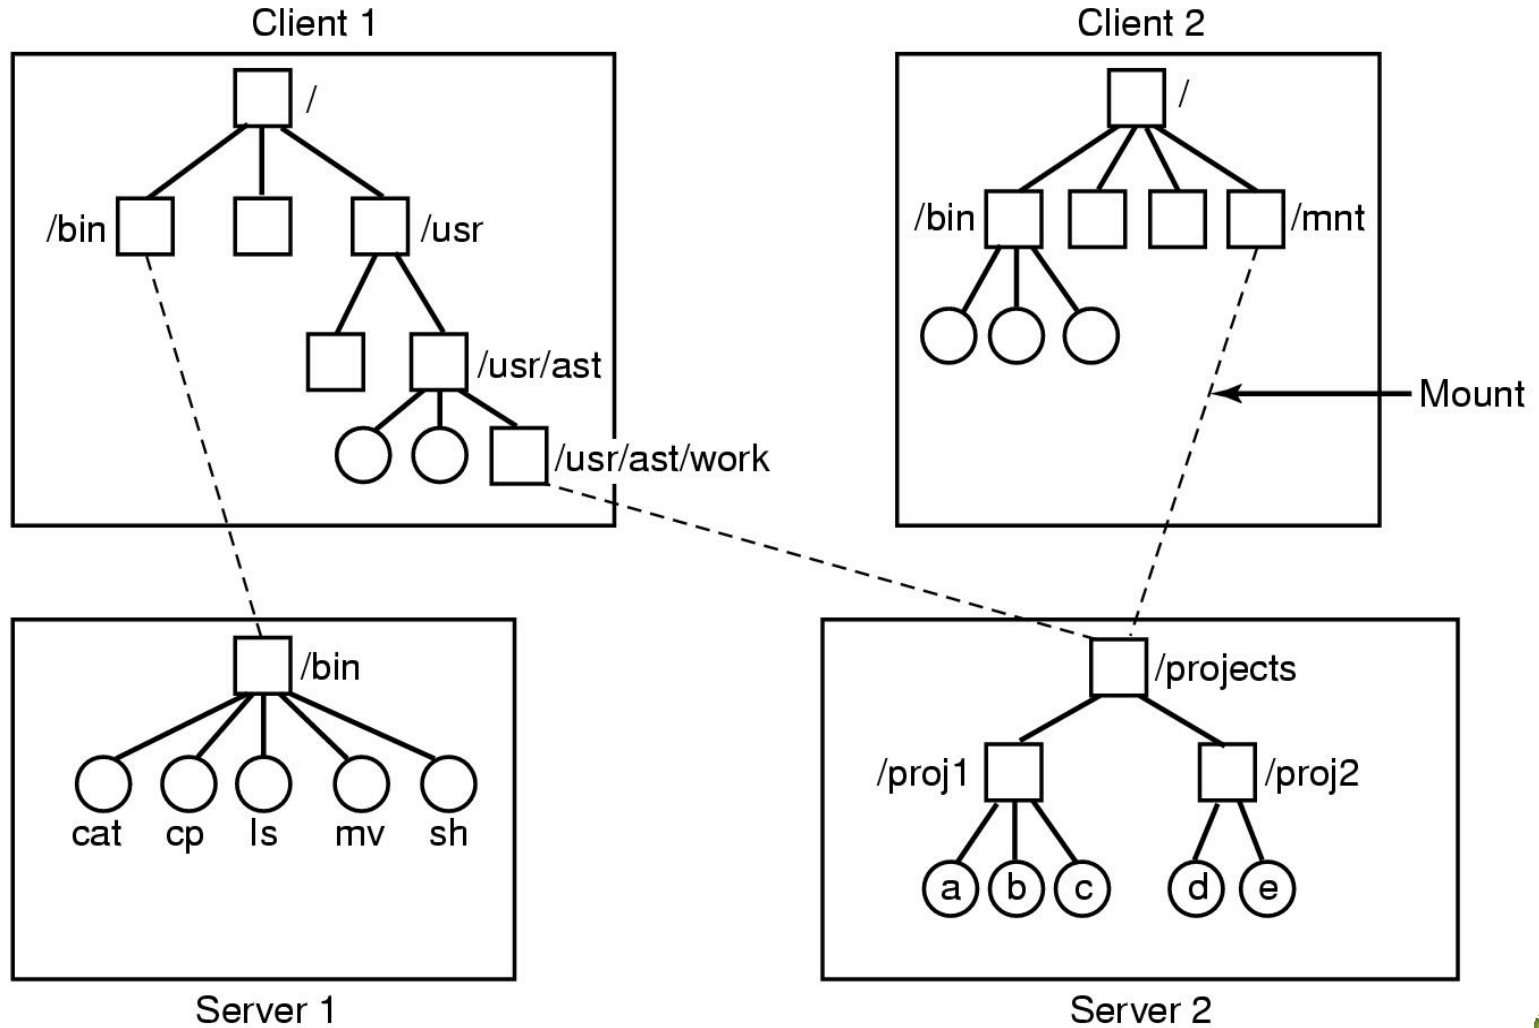
Network File System (NFS)

- ◆ Developed by Sun Microsystems in 1984
 - Used to join FSes on multiple computers as one logical whole
- ◆ Used commonly today with UNIX systems
- ◆ Assumptions
 - Allows arbitrary collection of users to share a file system
 - Clients and servers might be on different LANs
 - Machines can be clients and servers at the same time
- ◆ Architecture:
 - A server exports one or more of its directories to remote clients
 - Clients access exported directories by mounting them
 - ◆ The contents are then accessed as if they were local






Example






Three Independent File Systems



Mounting in NFS



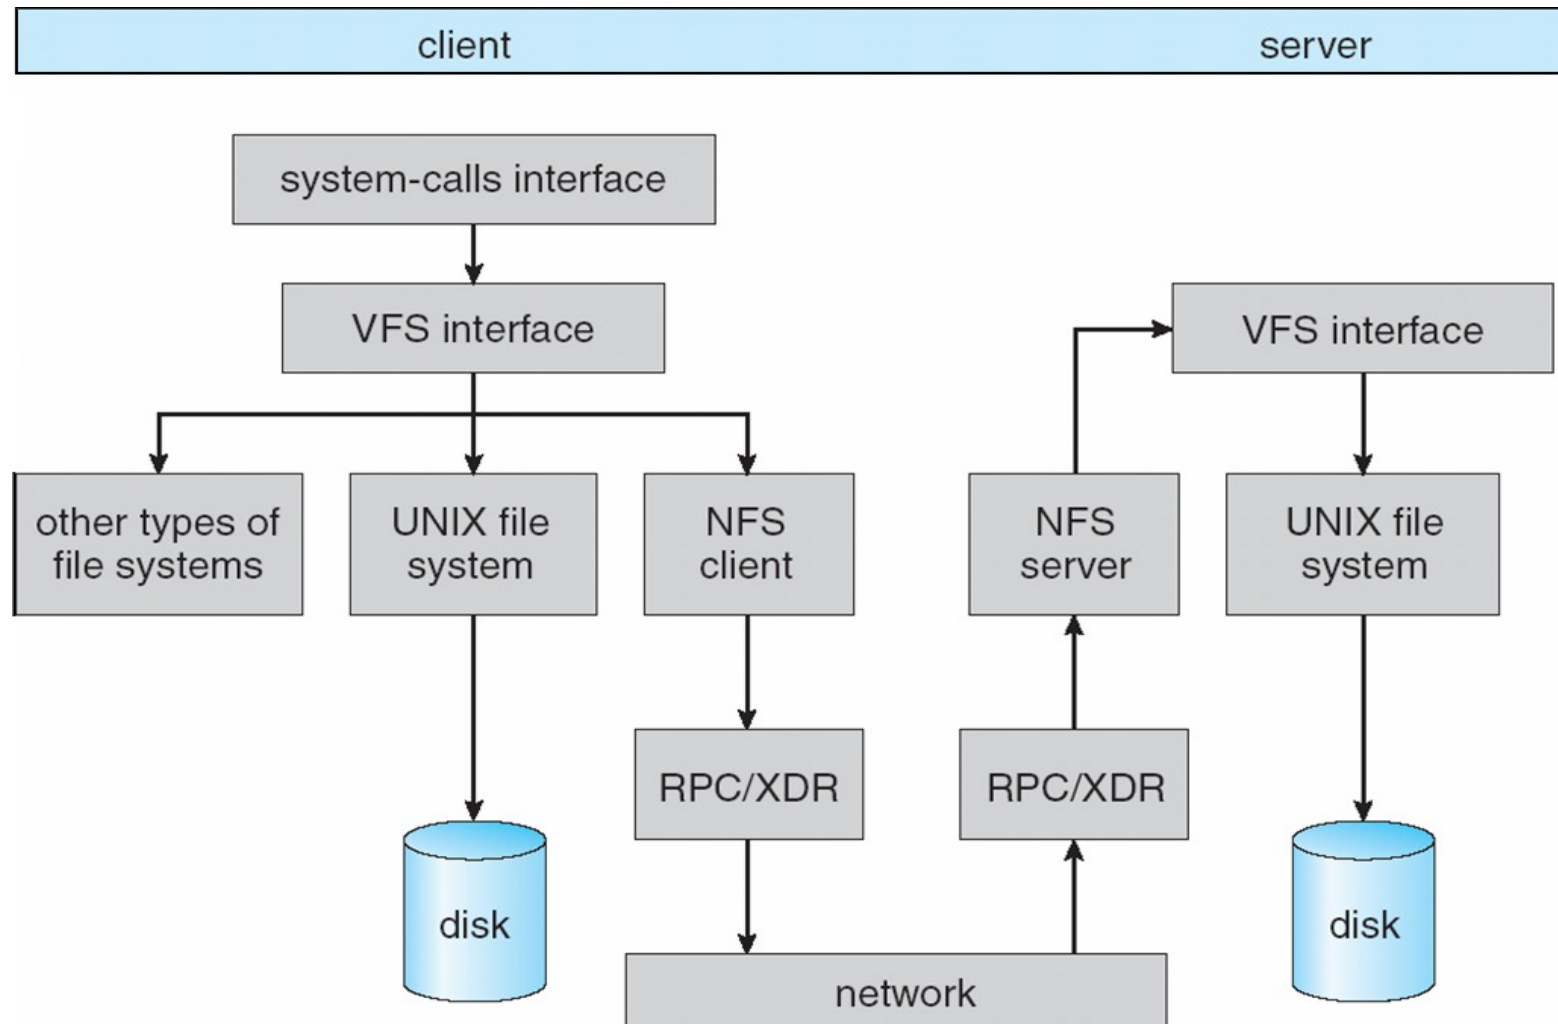
(a)
Mounts

(b)
Cascading
mounts





Schematic View of NFS Architecture





assignment

□ 11.3

□ 11.7

□ 11.8



End of Chapter 11

